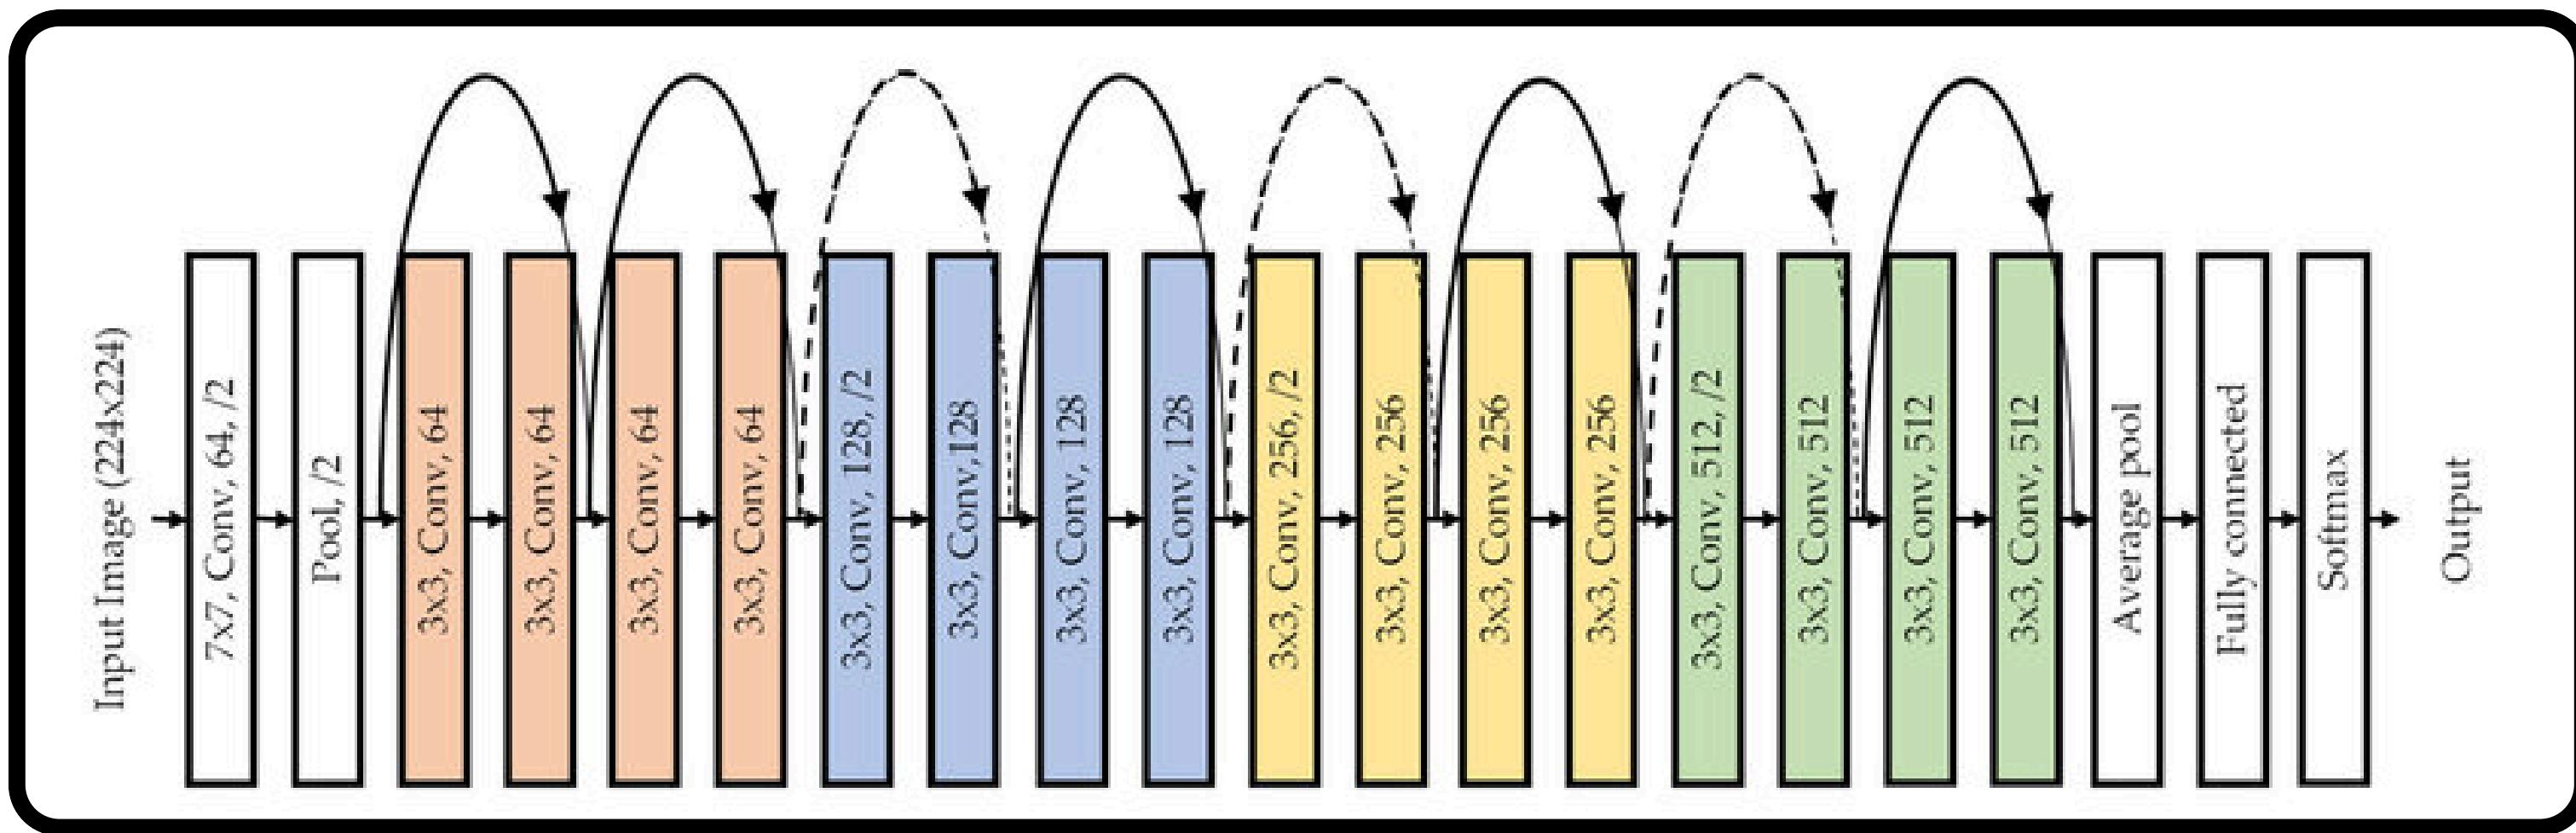


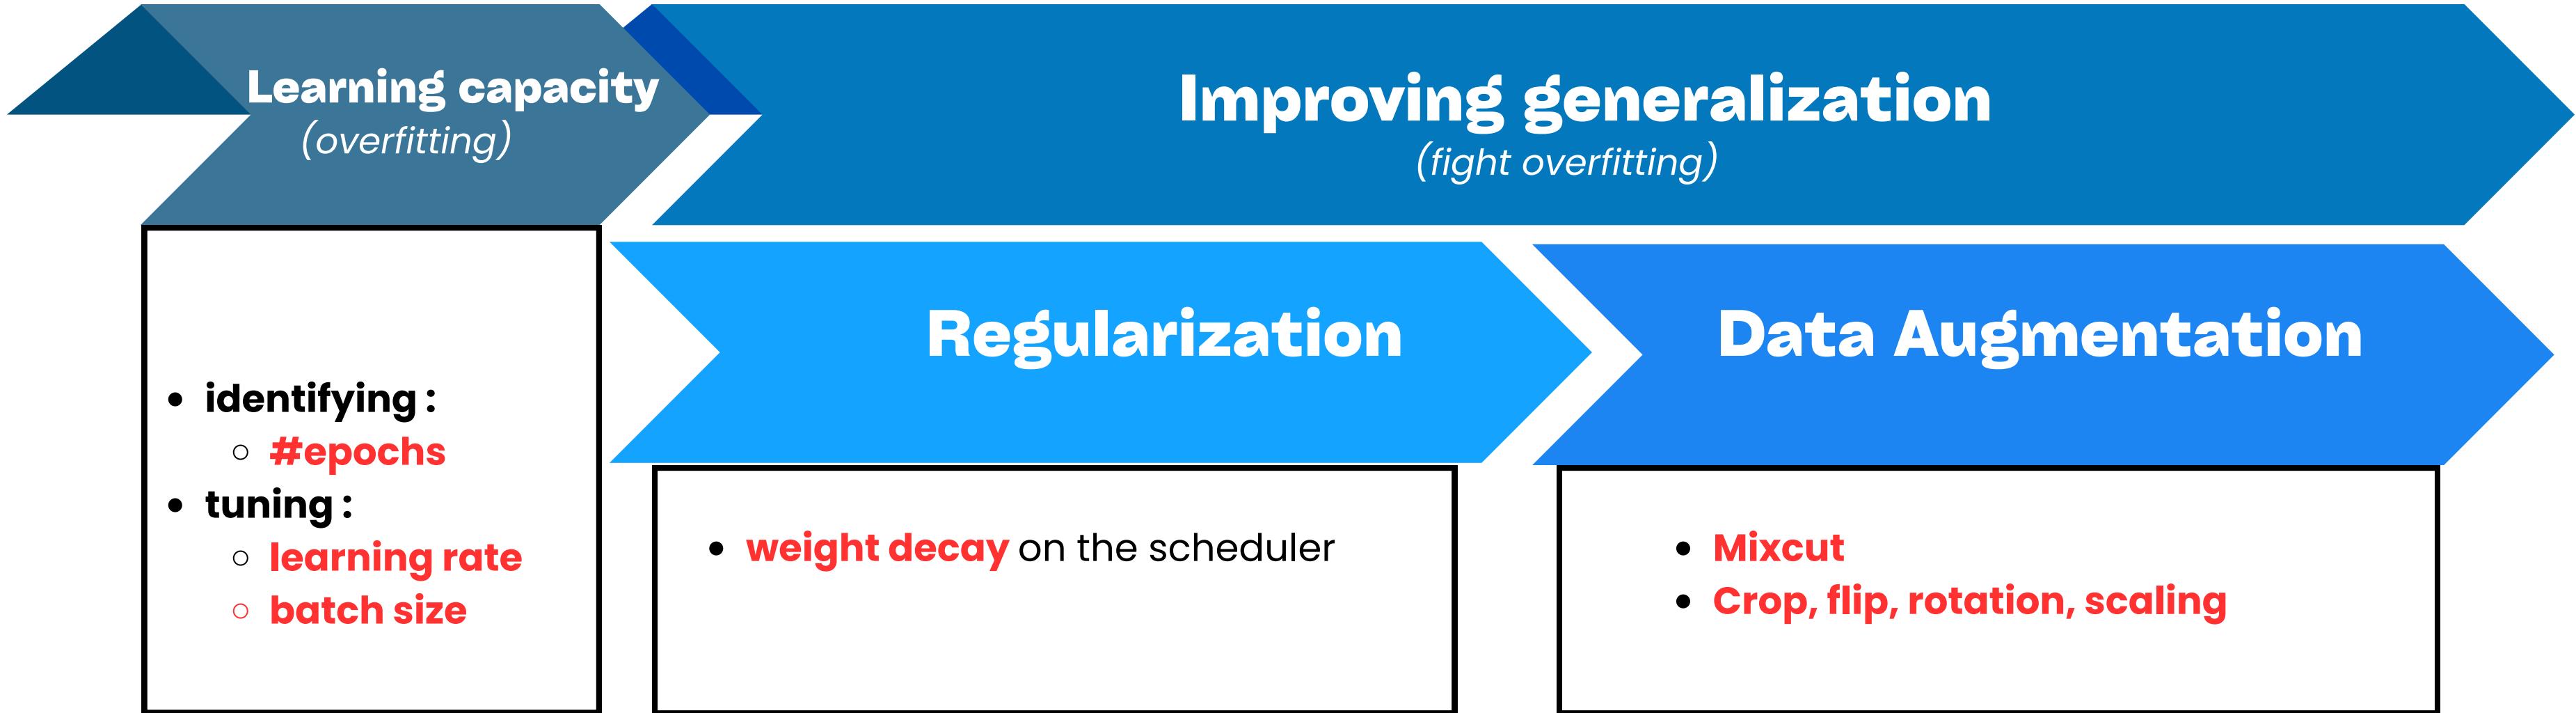
ResNet

(He, Zhang, Ren, & Sun, 2015)



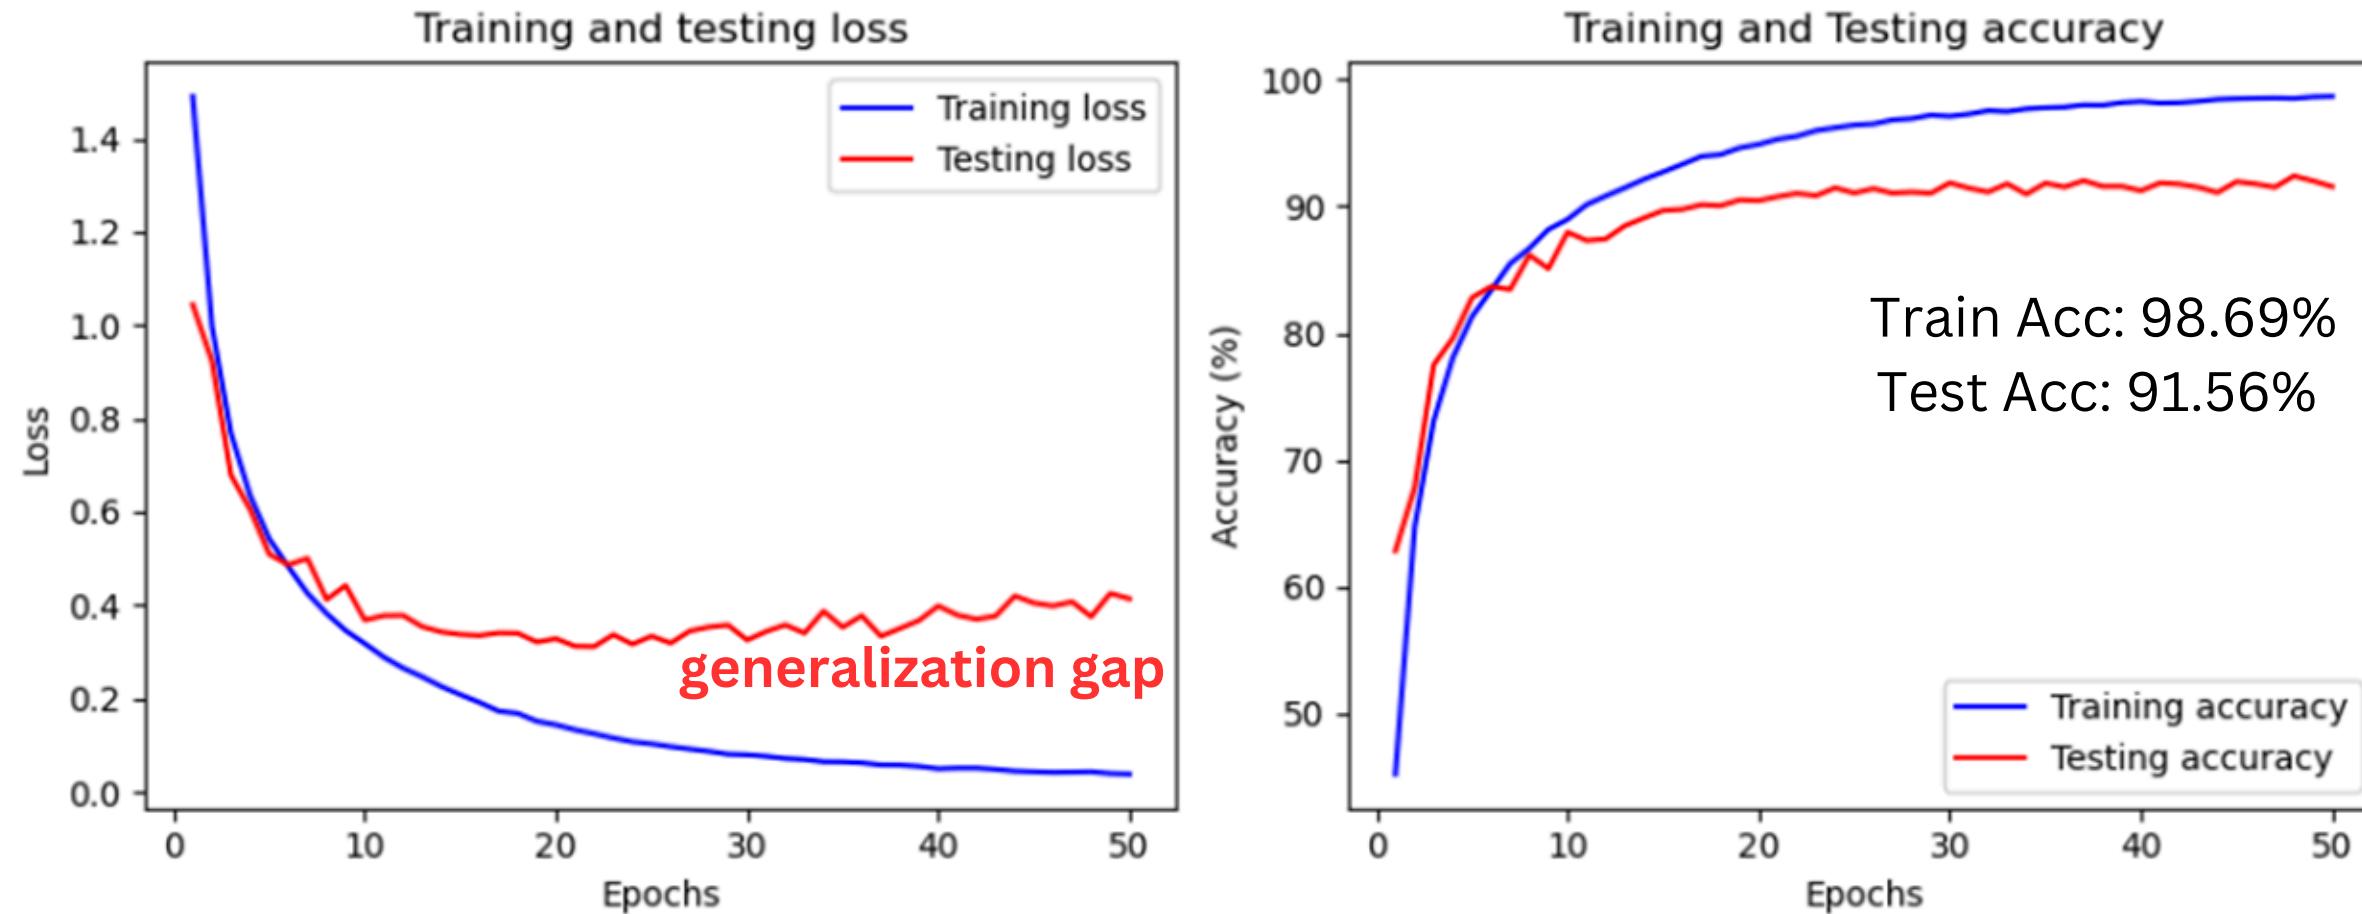
Xueyun WENG & Léa MIQUEU

Workflow



Learning capacity

Parameter	Value
# epochs	50
batch size	32
optimizer	Adam
learning rate	0.001



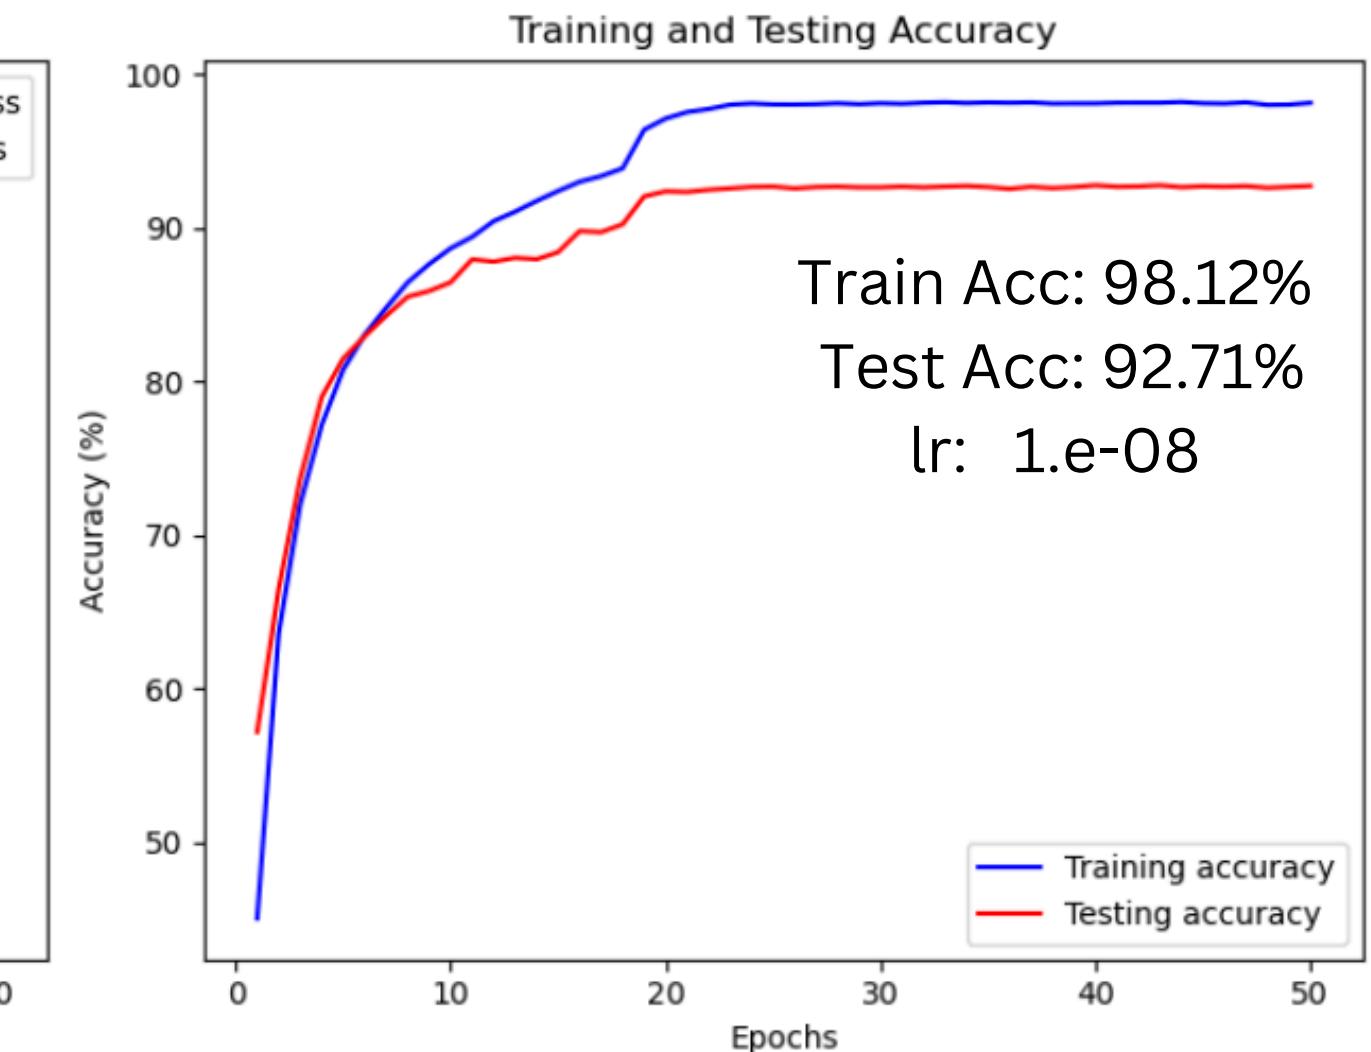
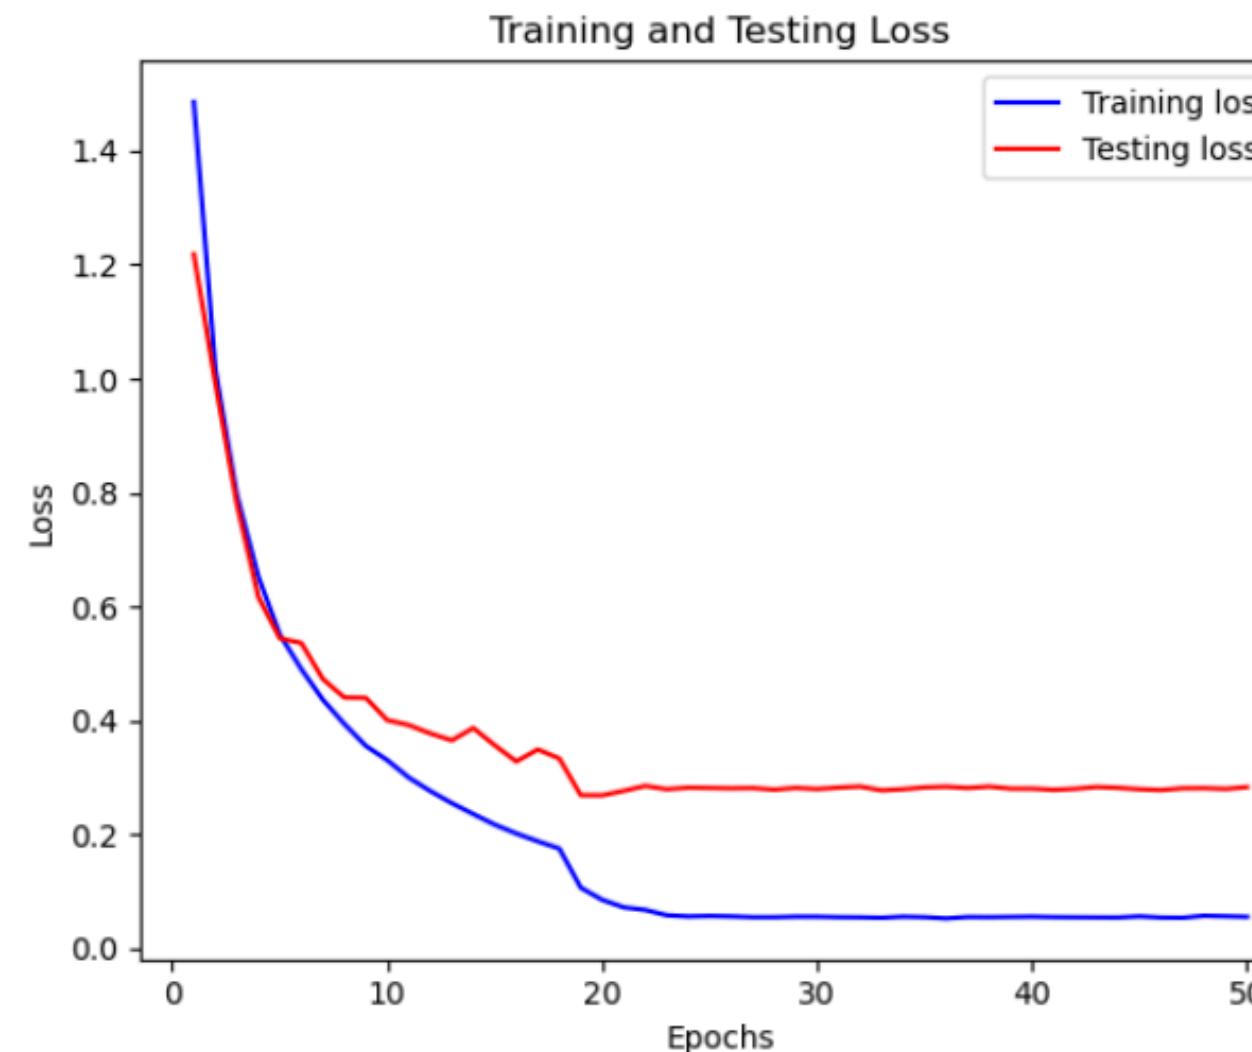
Observation:

✳️ the model overfits rapidly --> good learning capacity on training set but can't generalize

change the **batch size** for *Regularization Effect*

Learning capacity

Parameter	Value
# epochs	50
batch size	64
optimizer	Adam
learning rate	0.001
scheduler	On Plateau (0.10)

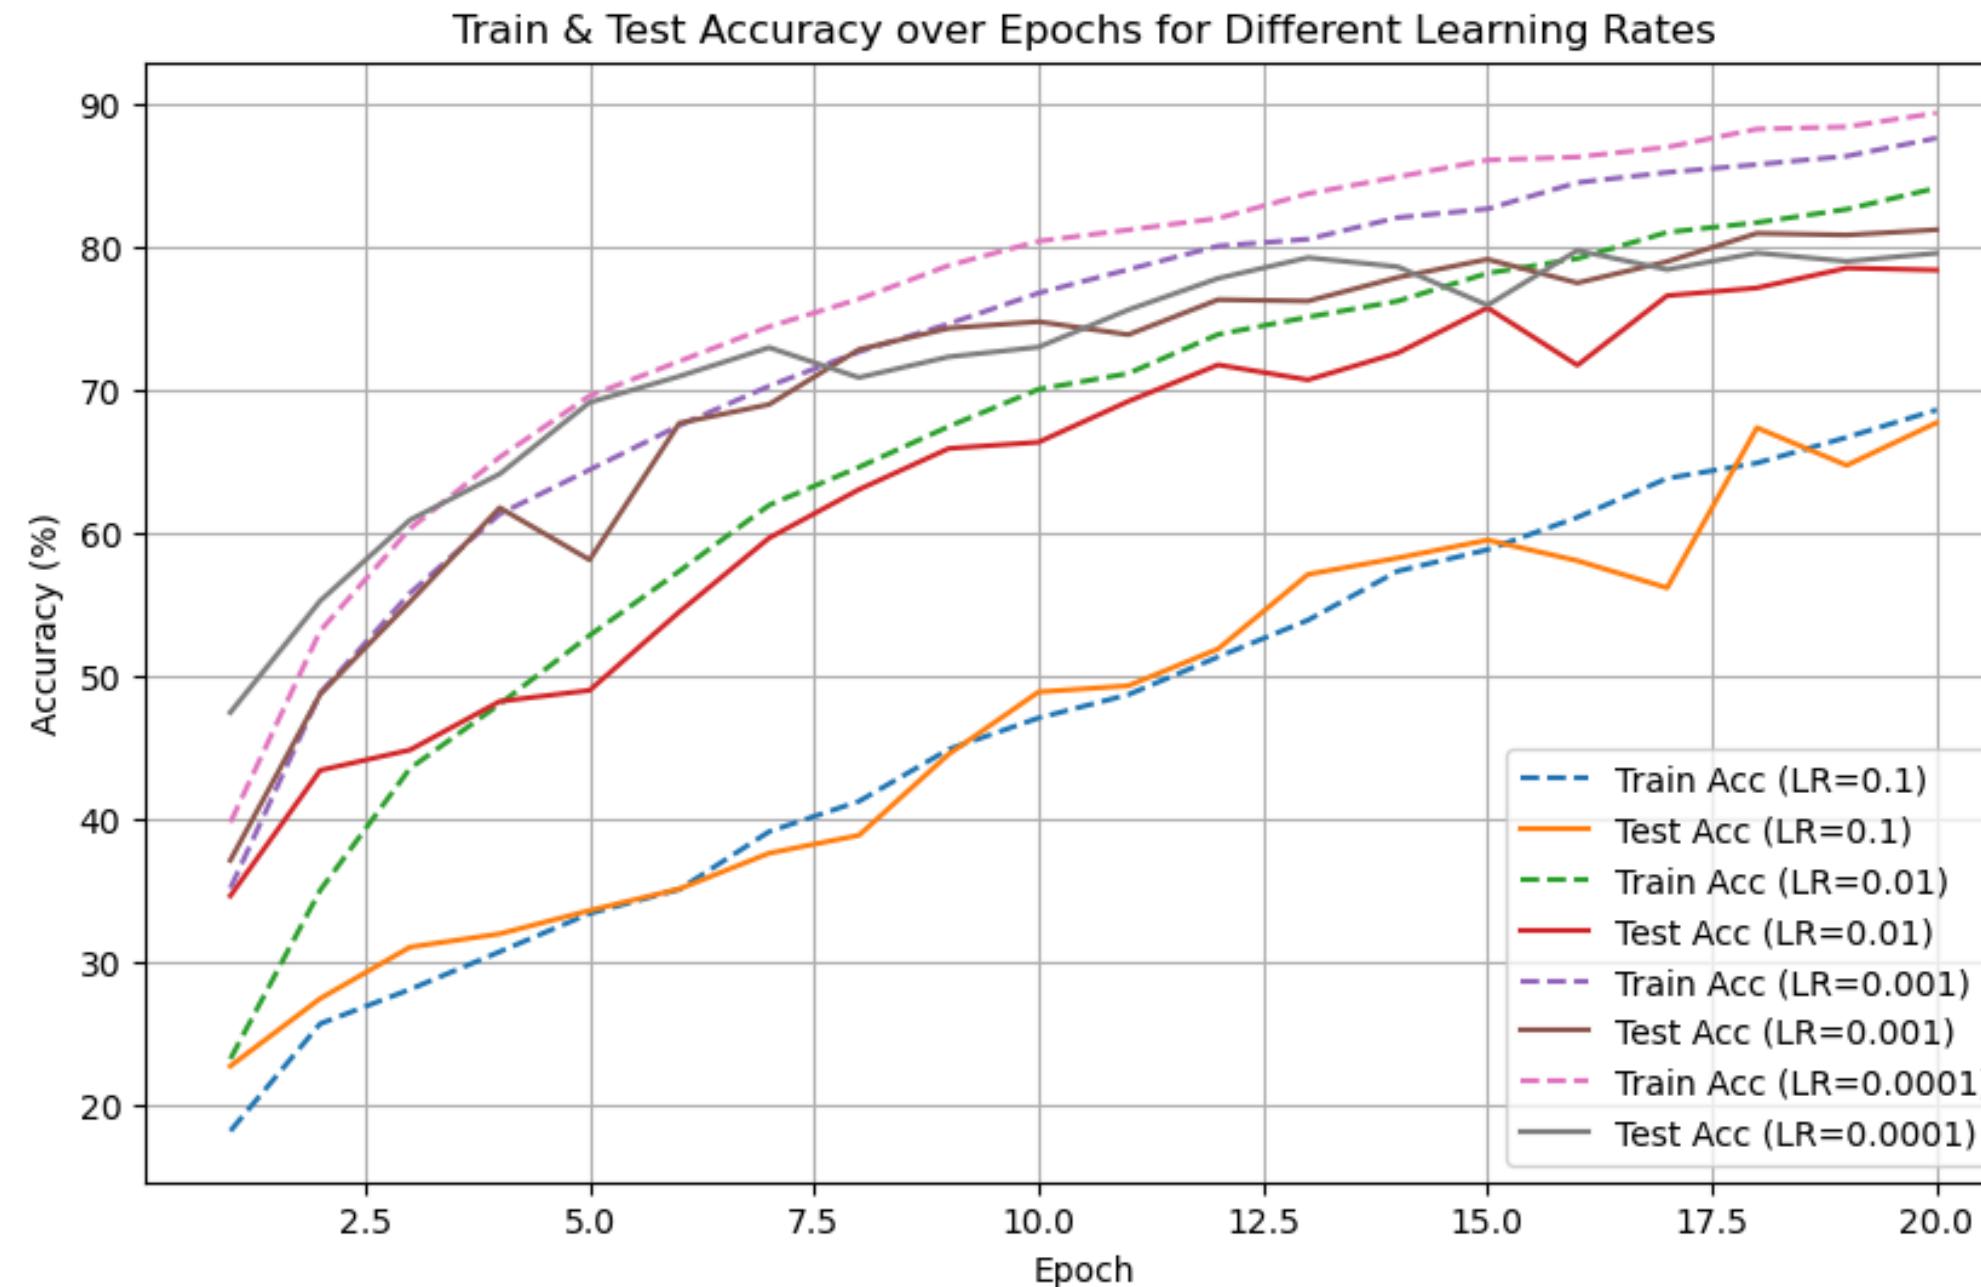


Observation:

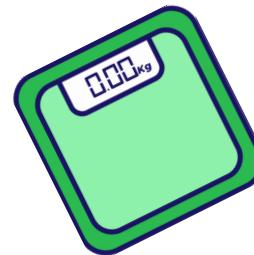
generalization (difference train/test : 7,13 --> 5.41)

Next: fight overfitting through model regularization and data augmentation

Starting Learning Rate



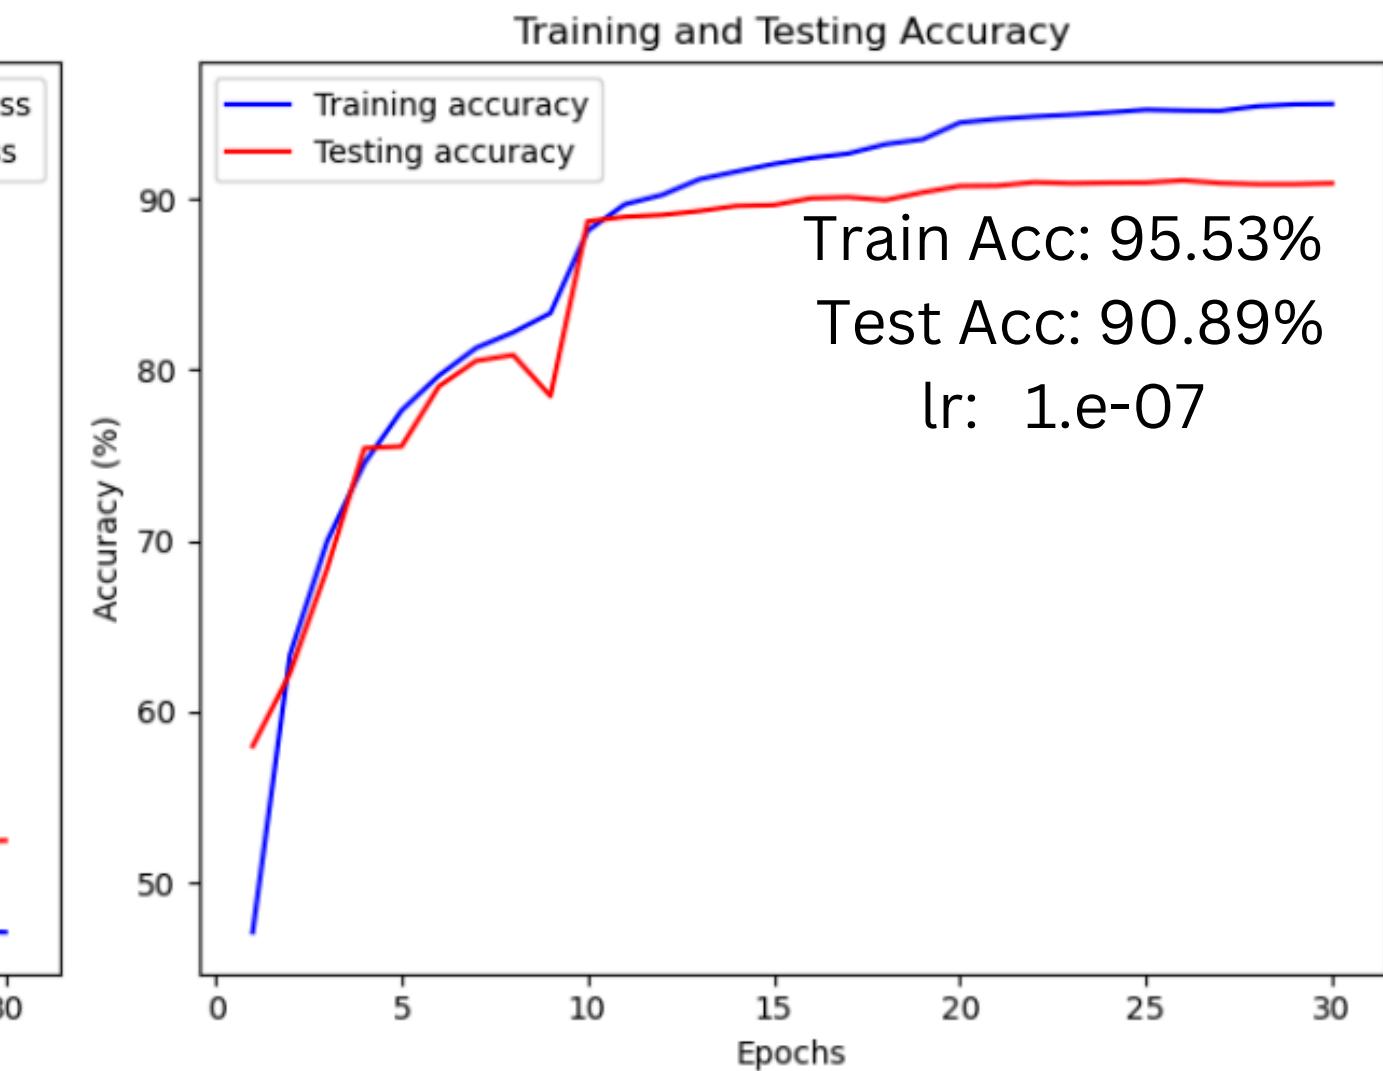
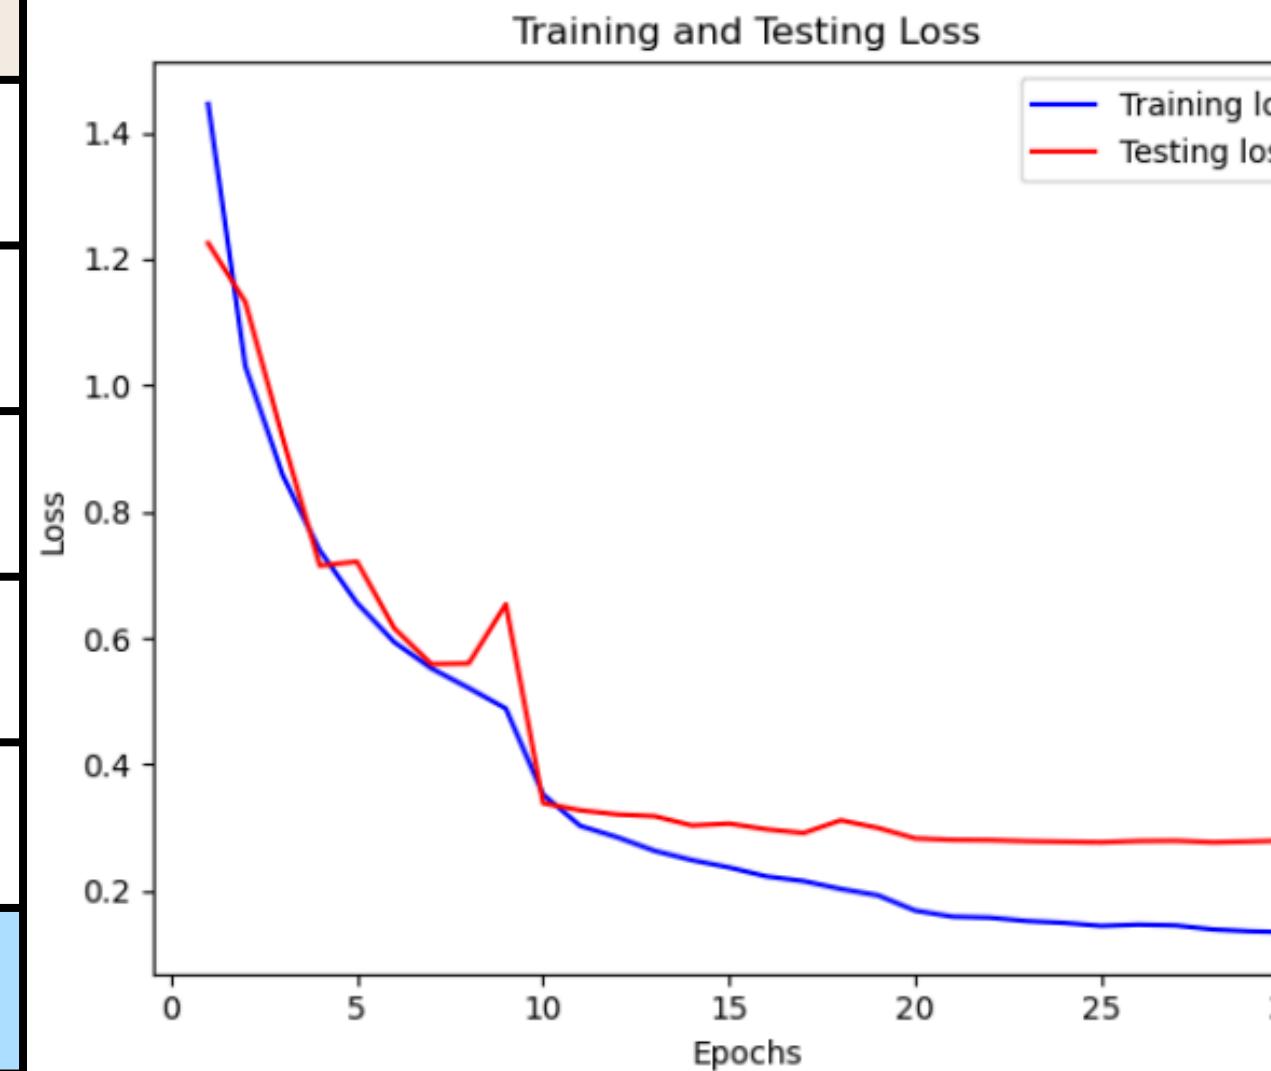
Parameter	Value
# epochs	50
batch size	64
optimizer	Adam
learning rate	0.001
scheduler	On Plateau (0.10)



Regularization with weight decay

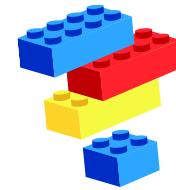
ℓ_2 penalty term is added to the loss, limits the growth of model weights

Parameter	Value
# epochs	30
batch size	64
optimizer	Adam
learning rate	0.001
scheduler	On Plateau (0.10)
weight decay	0.0005



Observation:

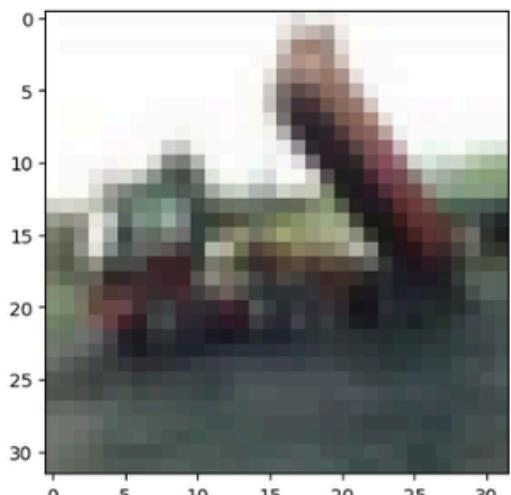
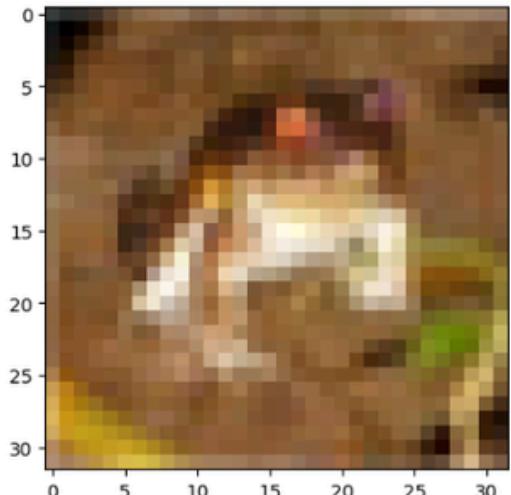
- slightly getting rid of overfitting



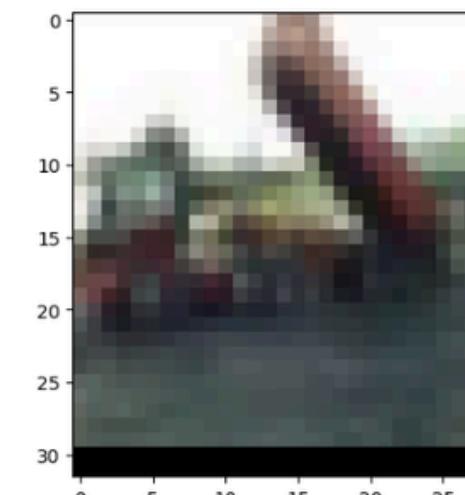
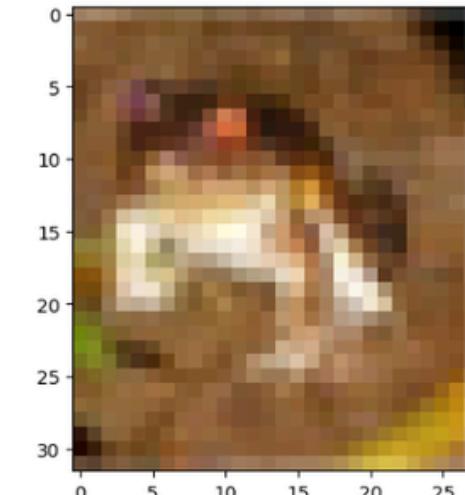
Data Augmentation

Help generalization by sampling training examples from a larger distribution using randomized transforms

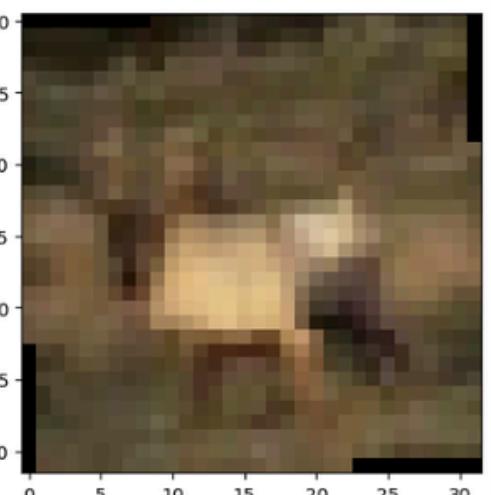
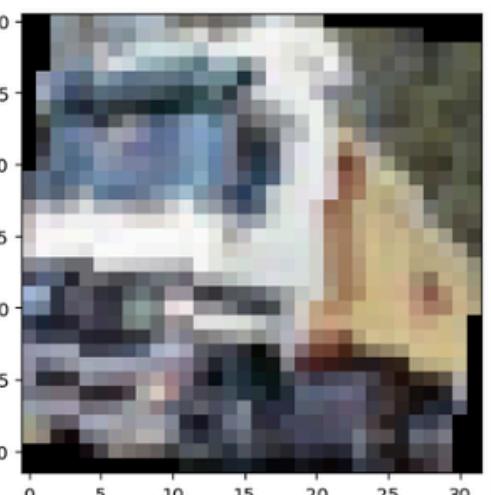
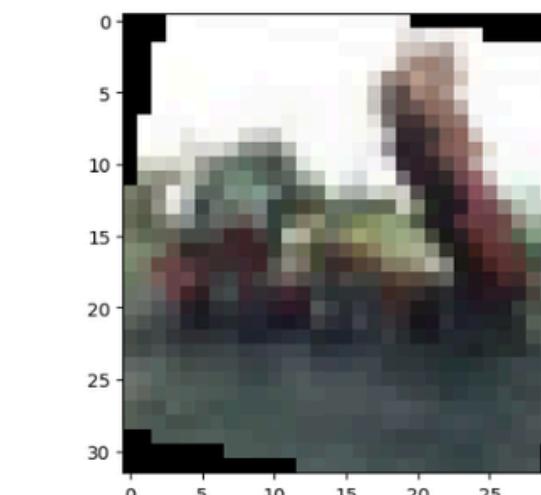
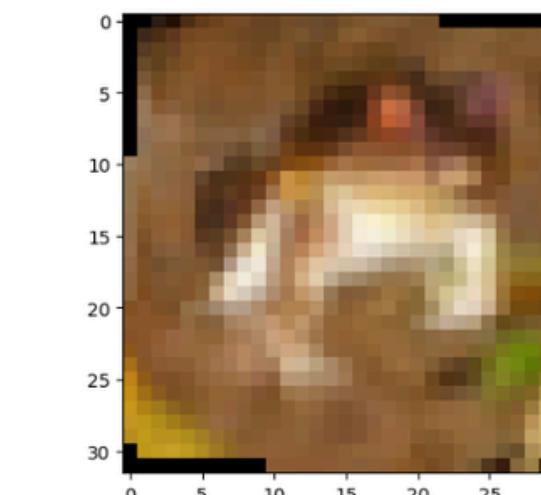
No Data Augmentation

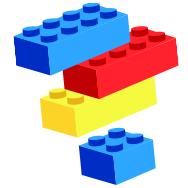


Crop and Flip



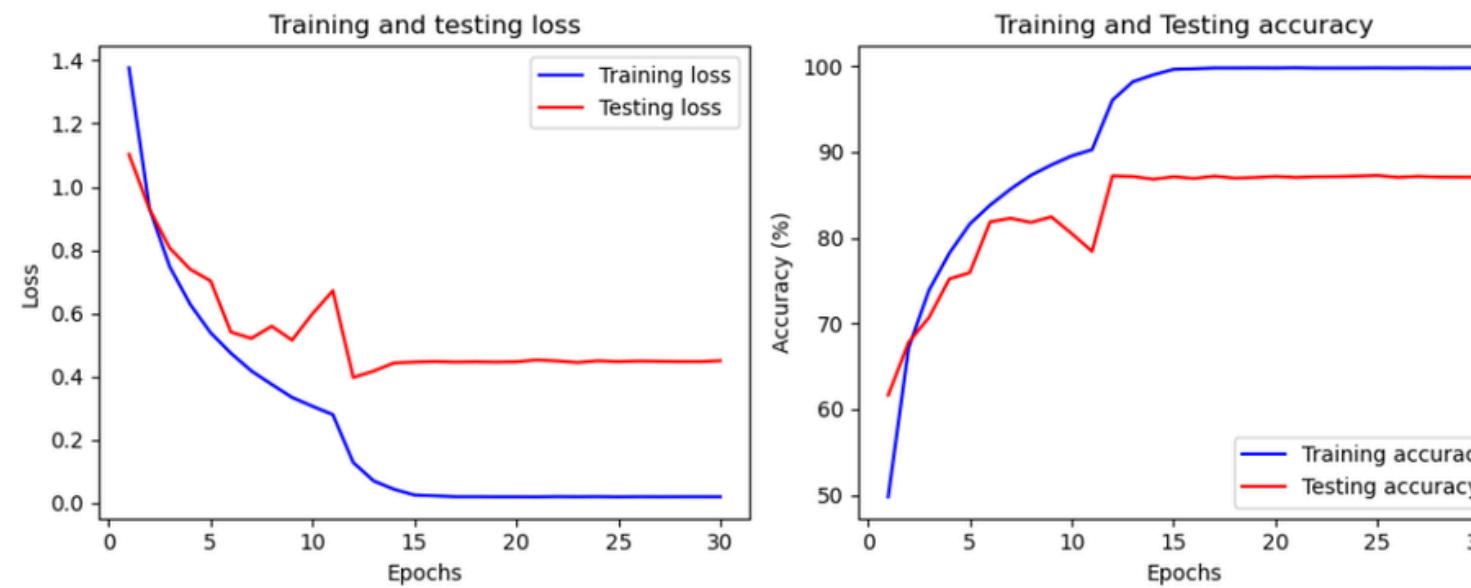
Scaling and Rotation



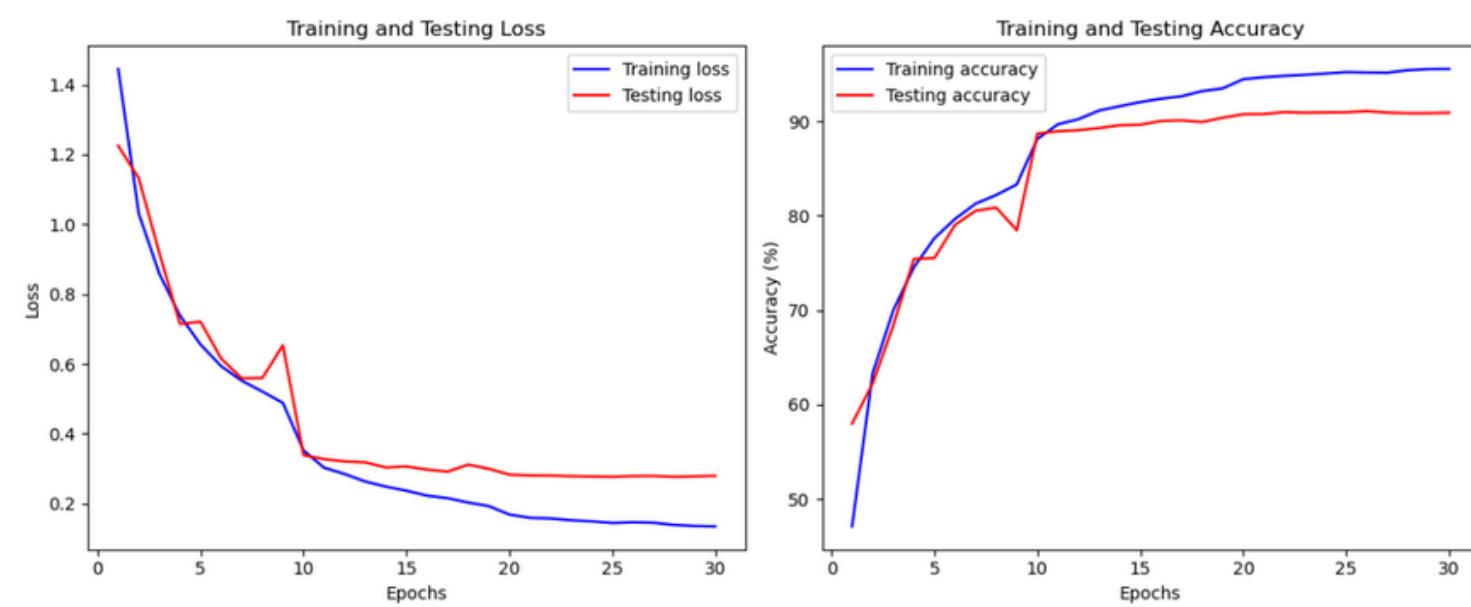


Data Augmentation

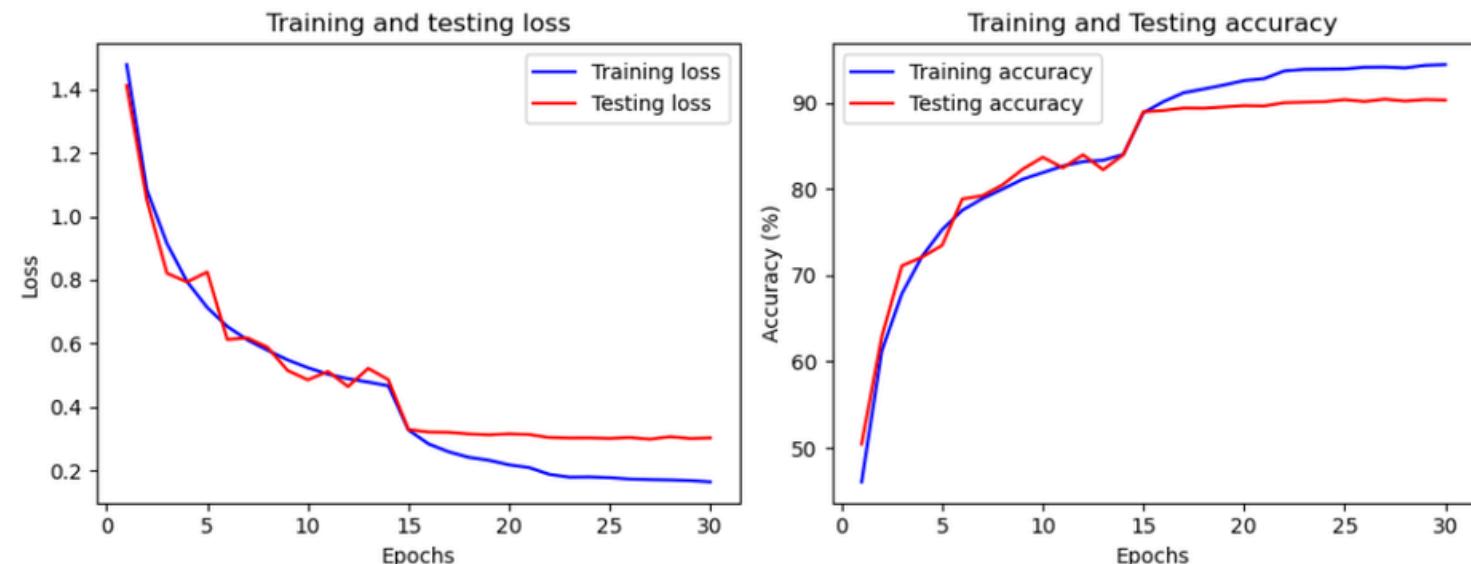
No Data Augmentation

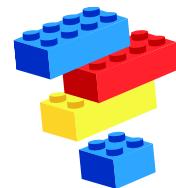


Crop and Flip

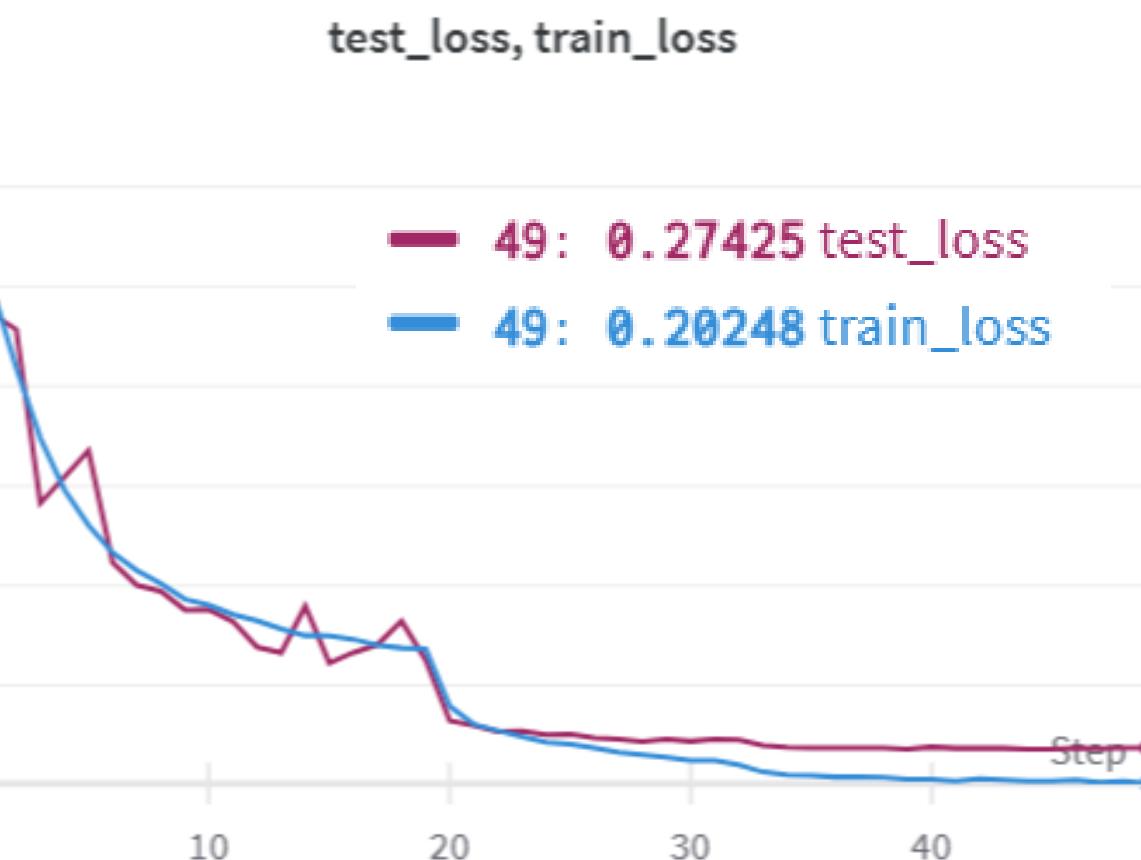
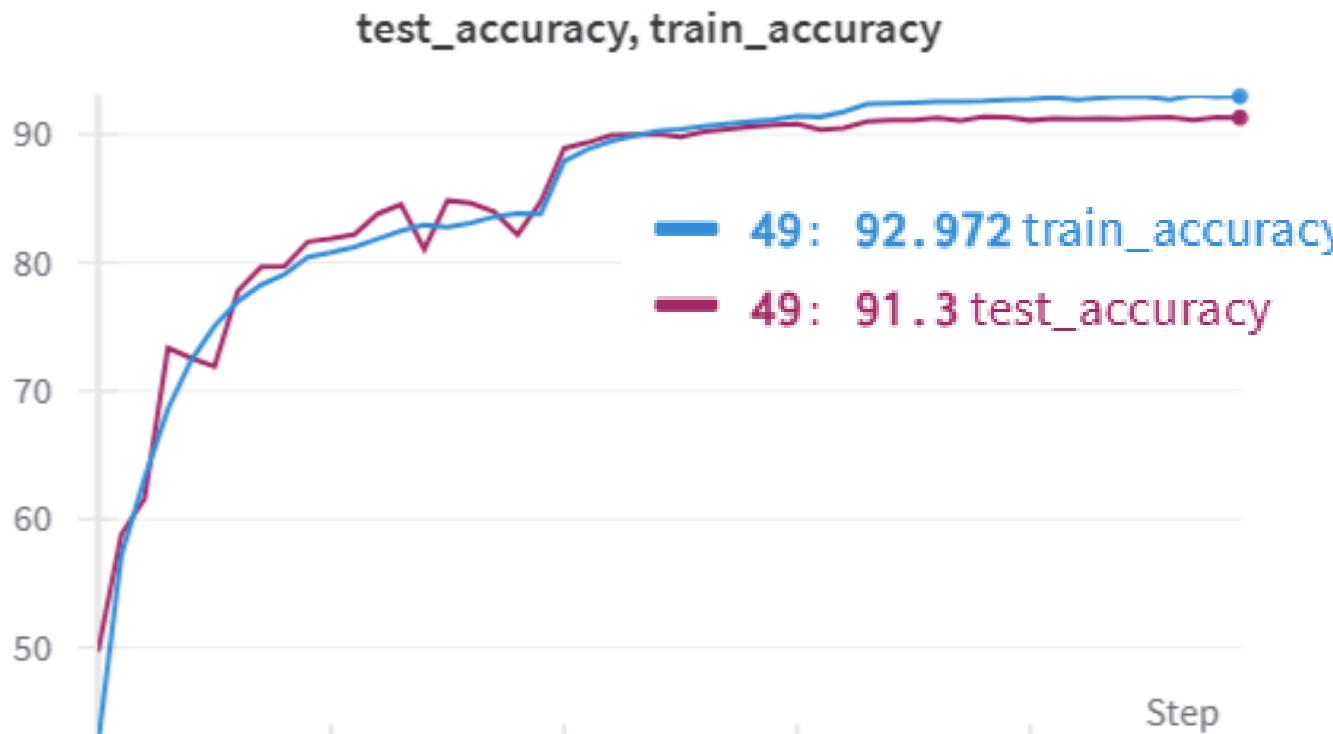


Scaling and Rotation



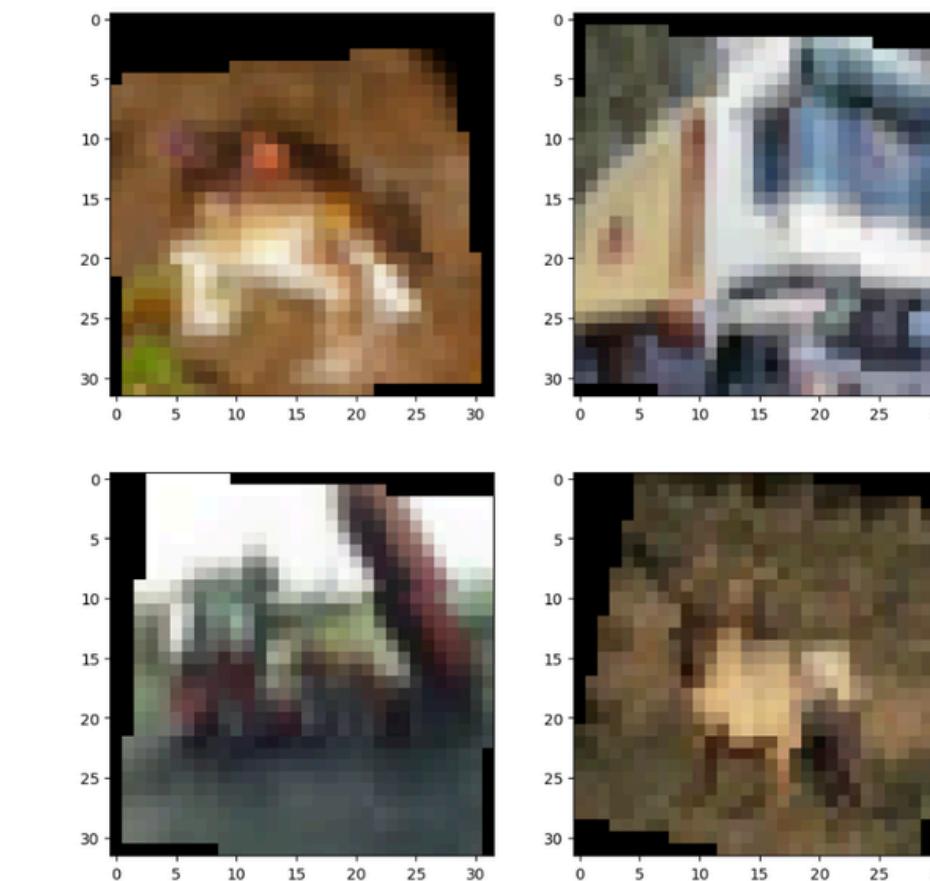


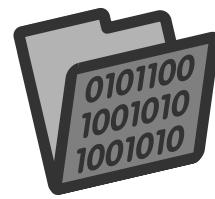
Data Augmentation



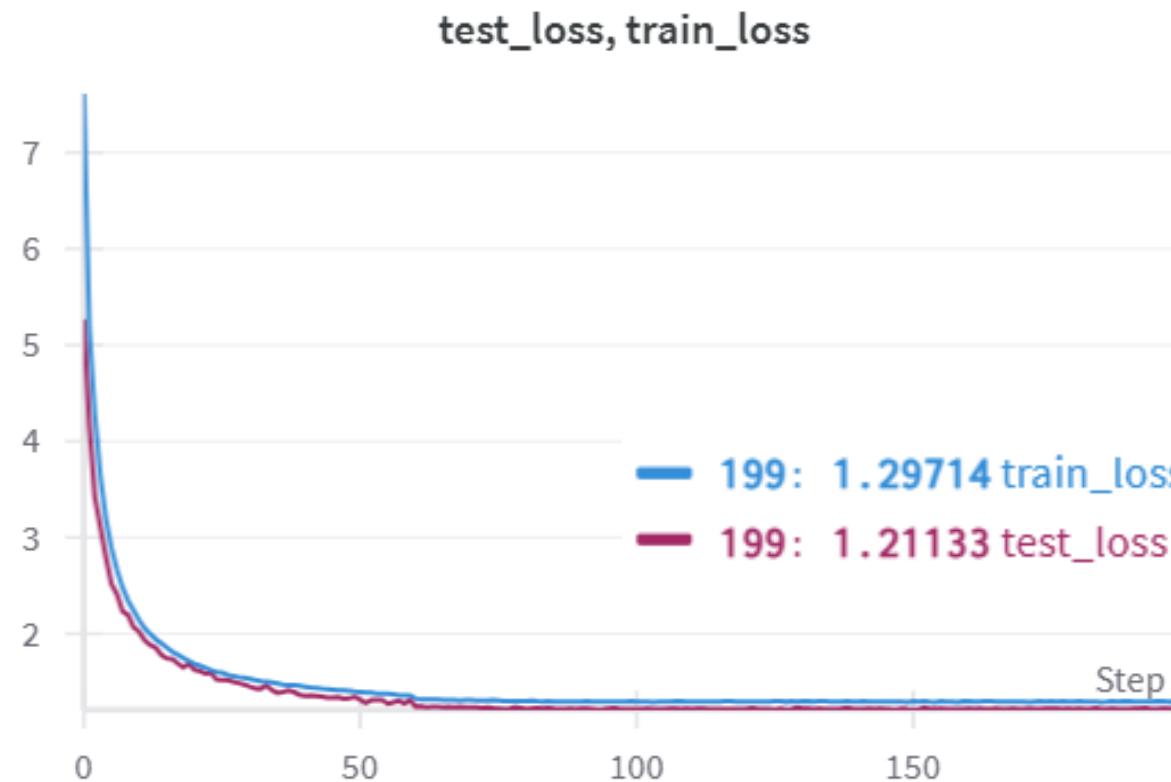
Crop + Flip + Scaling + Rotation

fancy

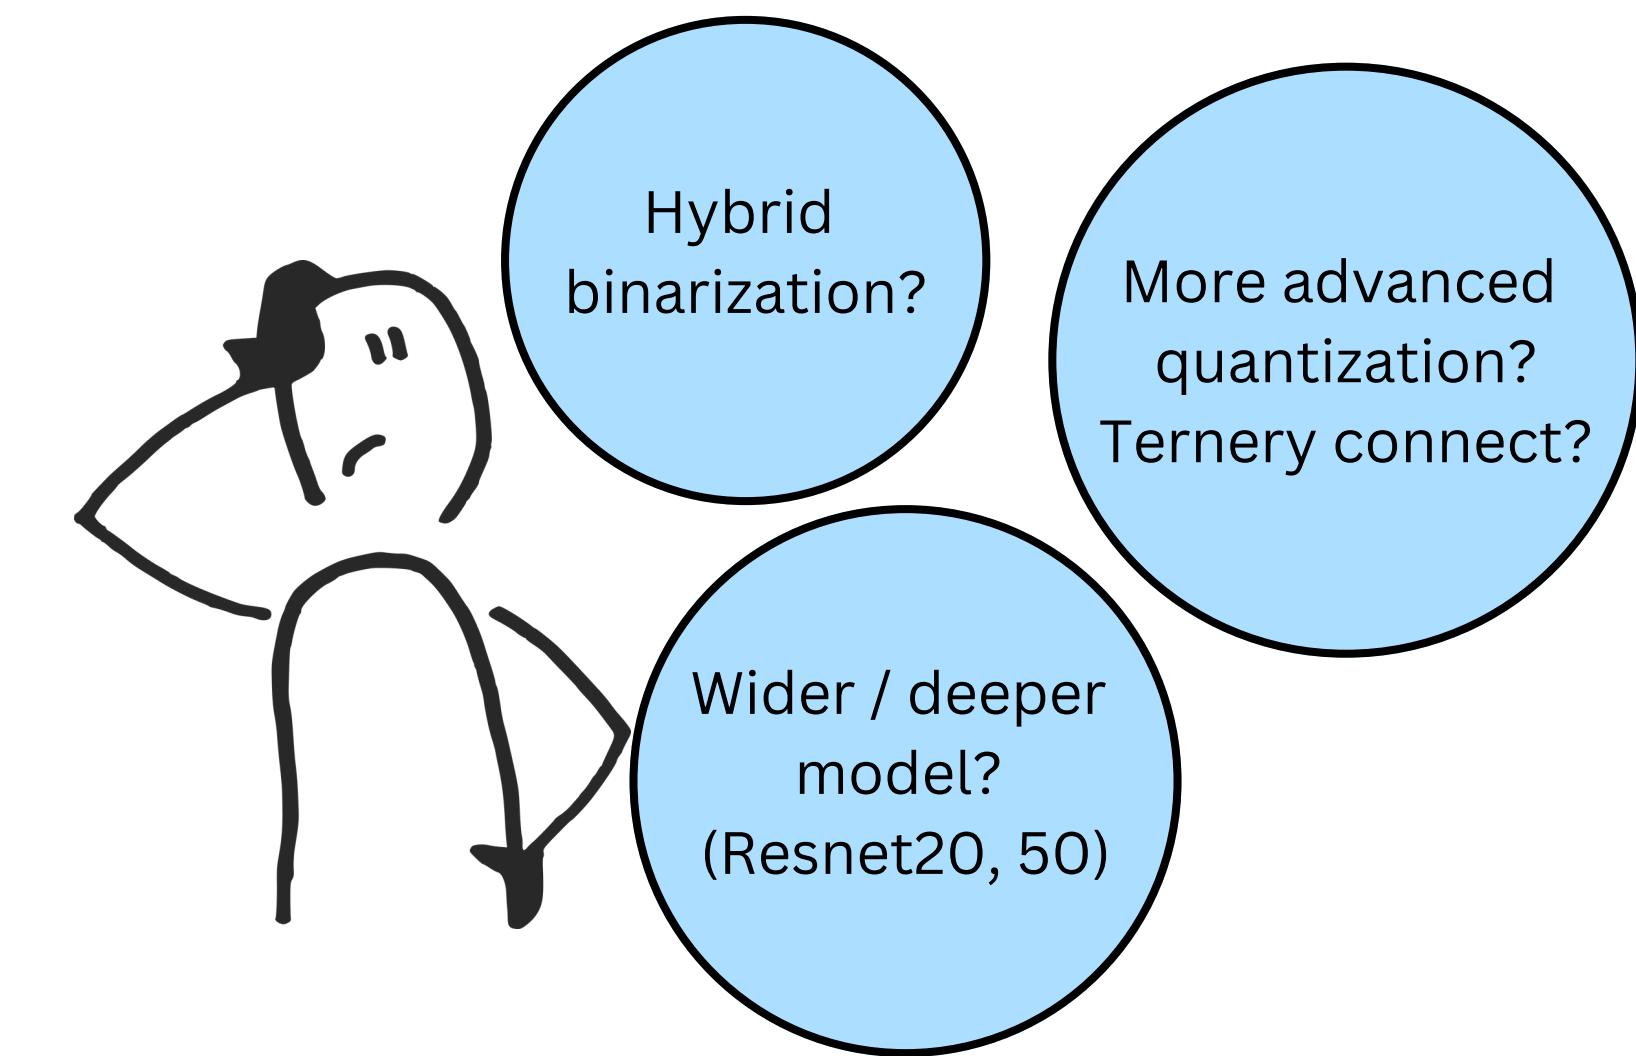




Quantization - Binary Connect

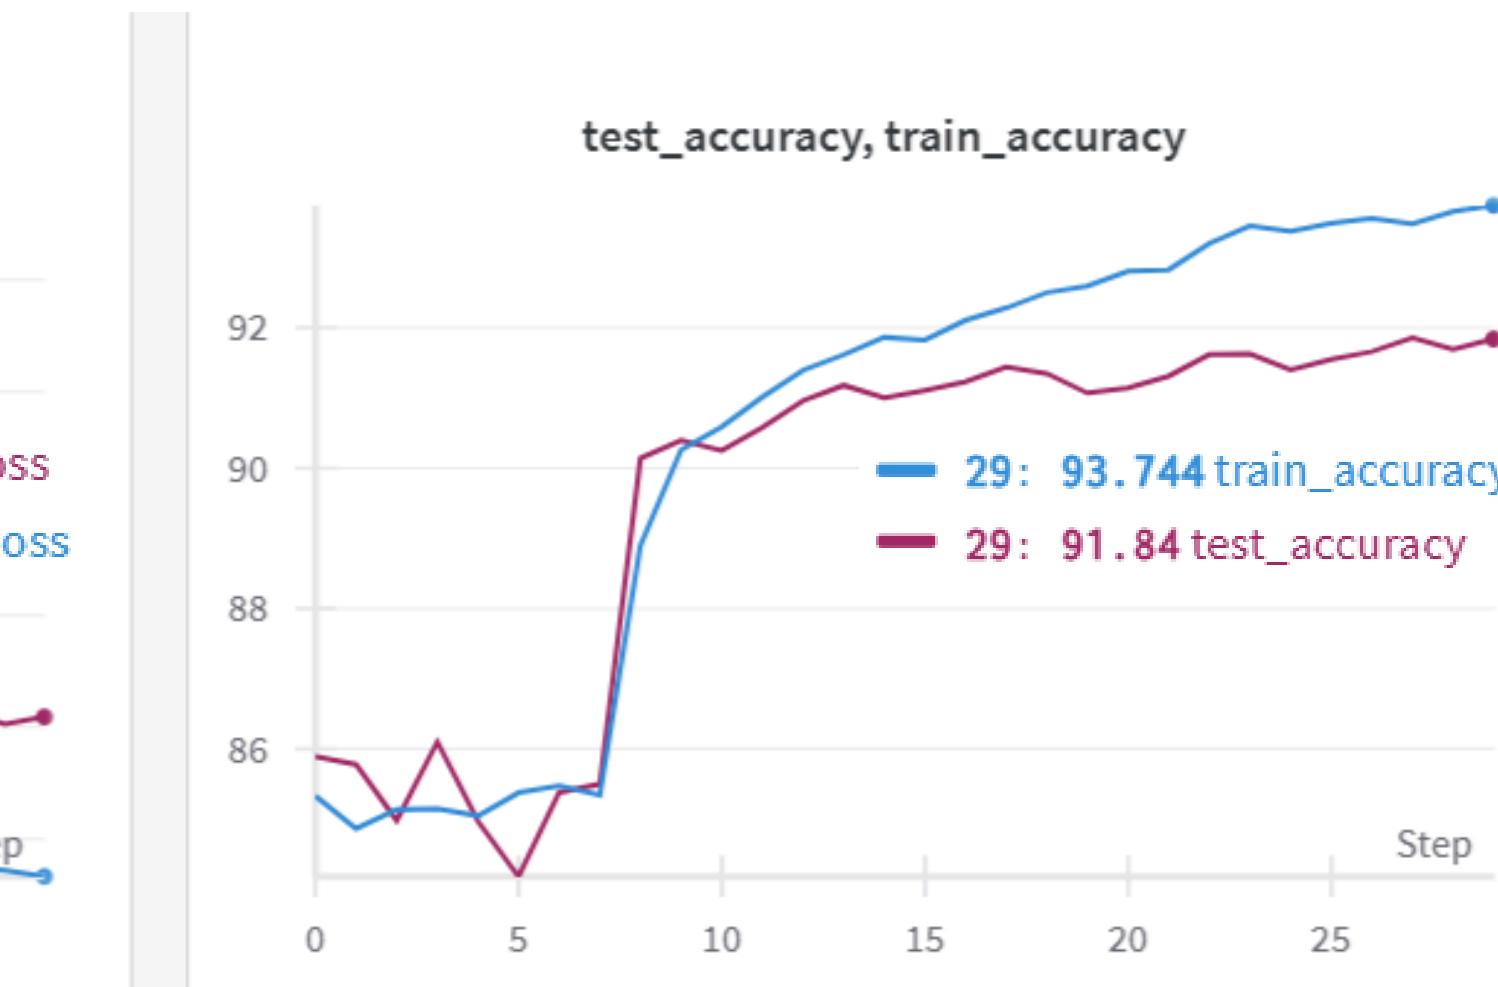


- Bad and saturating accuracies
- Losses stagnate



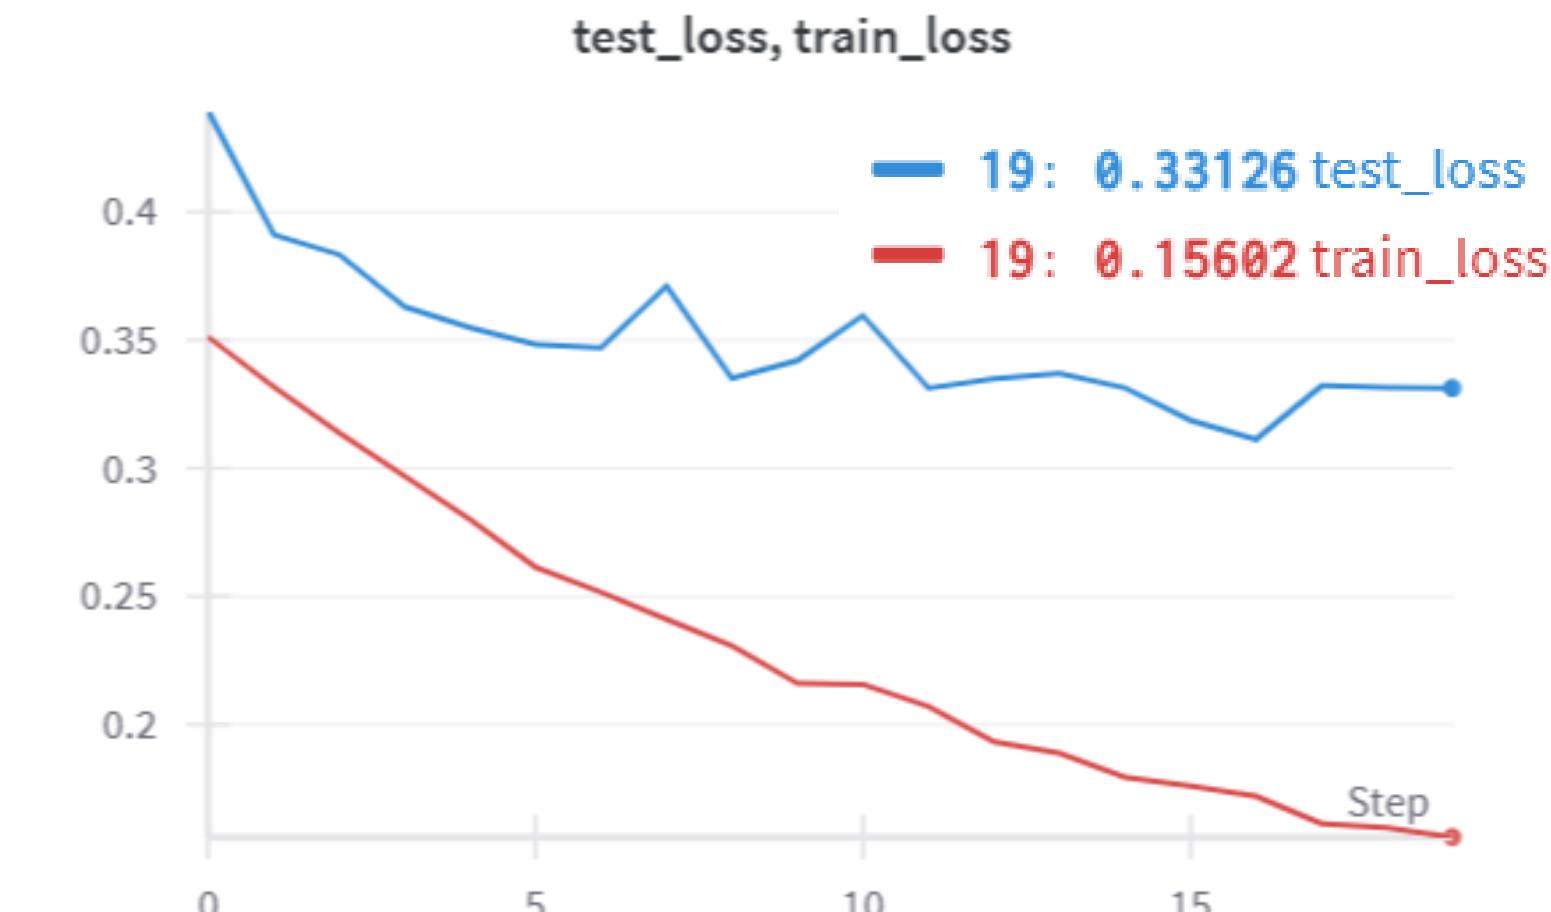
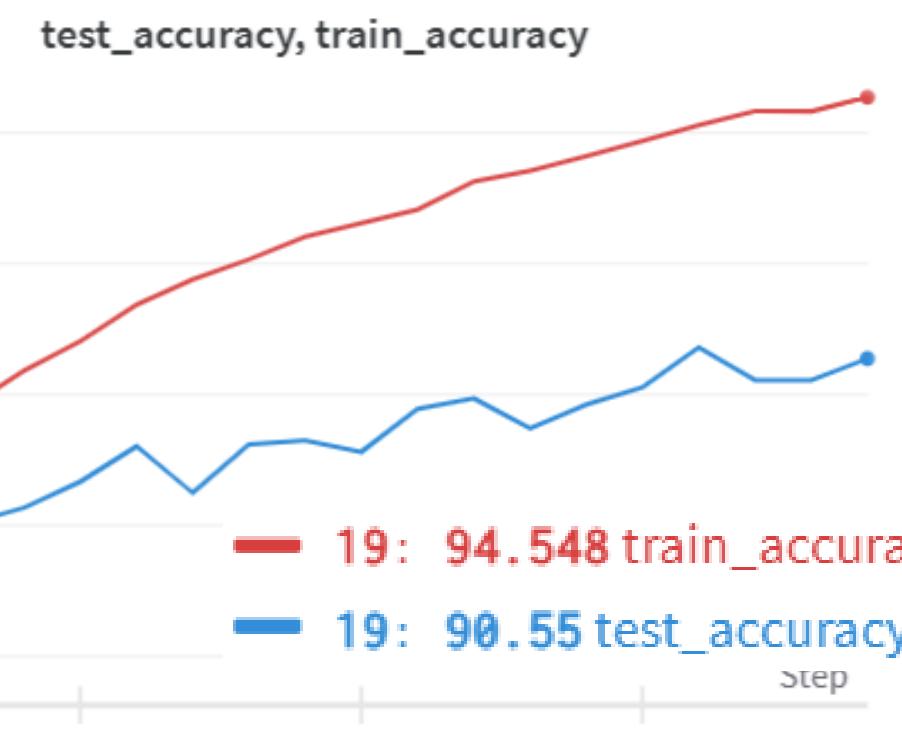


Simple pruning





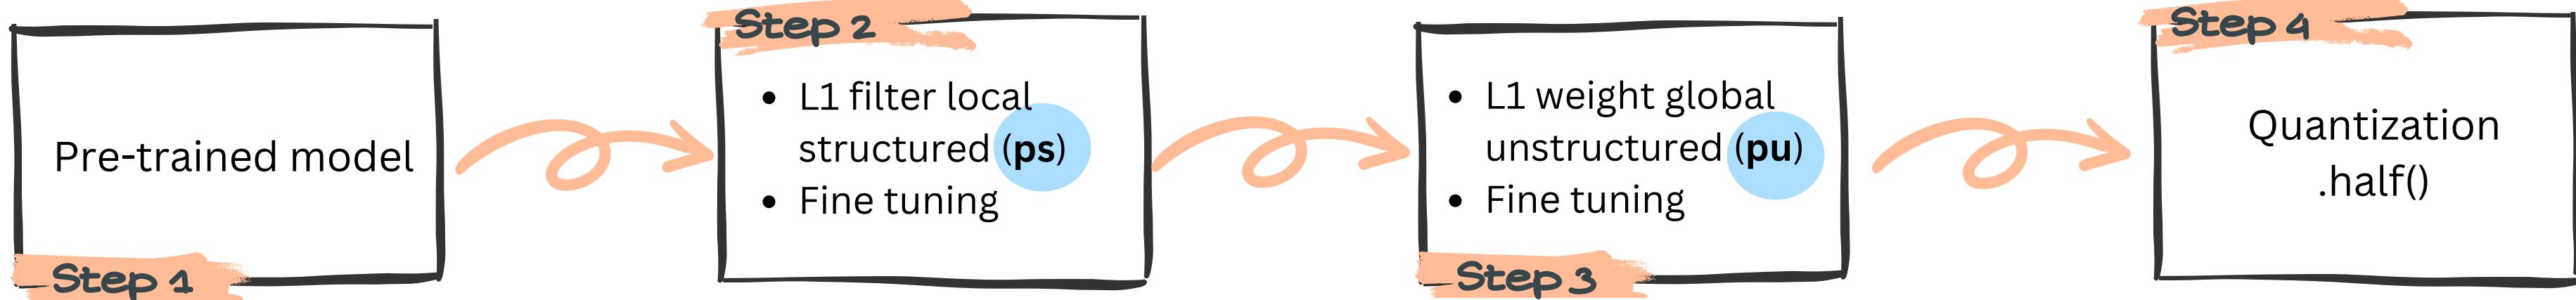
Simple pruning



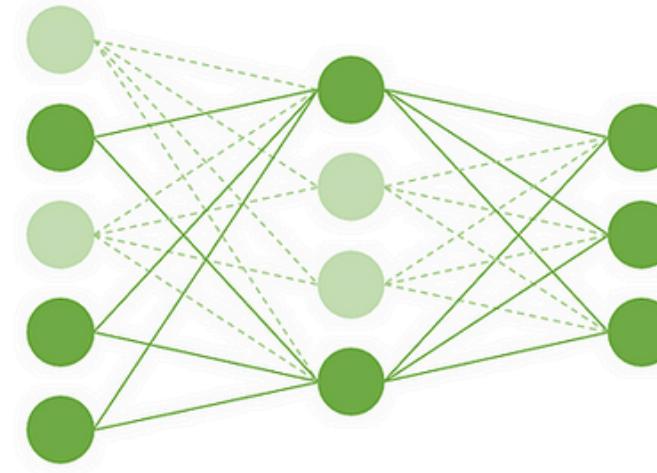
- higher sparsity --> lower accuracy
- needs more epoch to stabilize



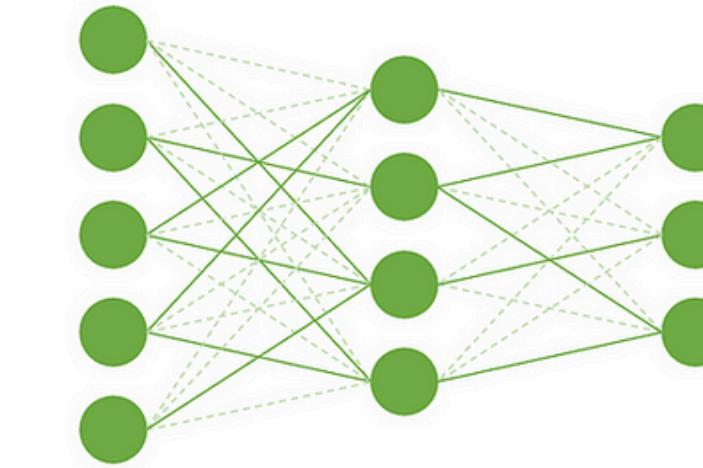
Pruning strategy



Structured Pruning



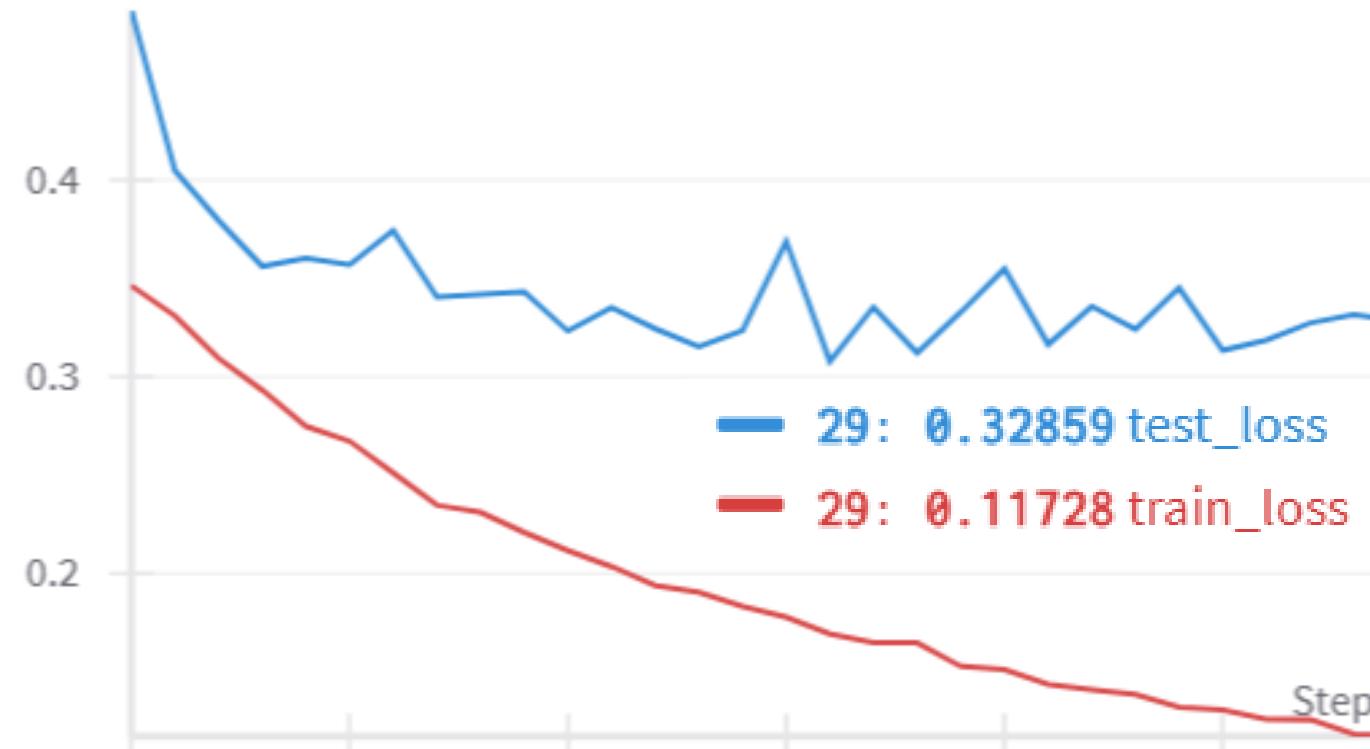
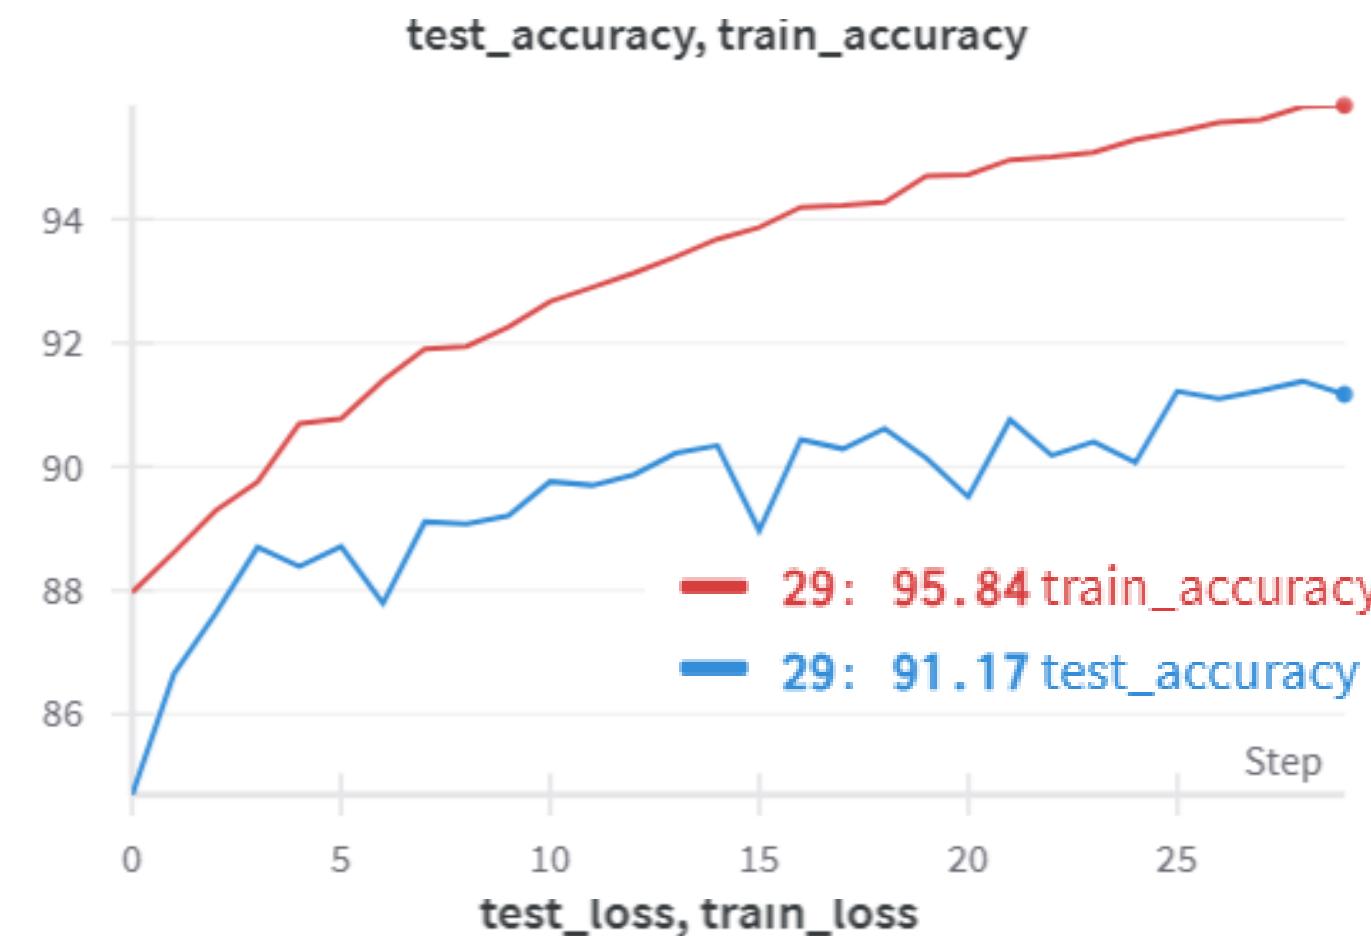
Unstructured Pruning



Start with structured pruning (e.g., remove 30-50% filters).
Fine-tune with unstructured pruning to reach higher sparsity.
Retrain after each pruning step to recover accuracy.



Pruning strategy

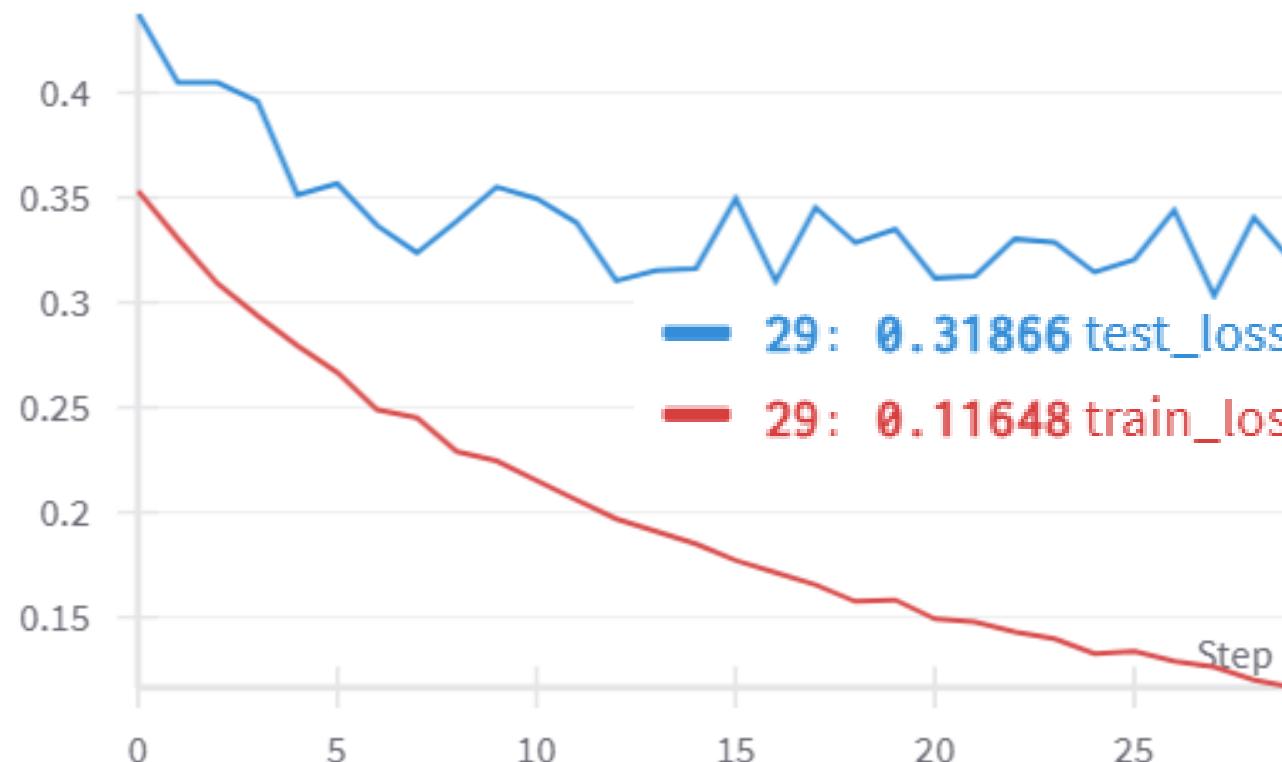
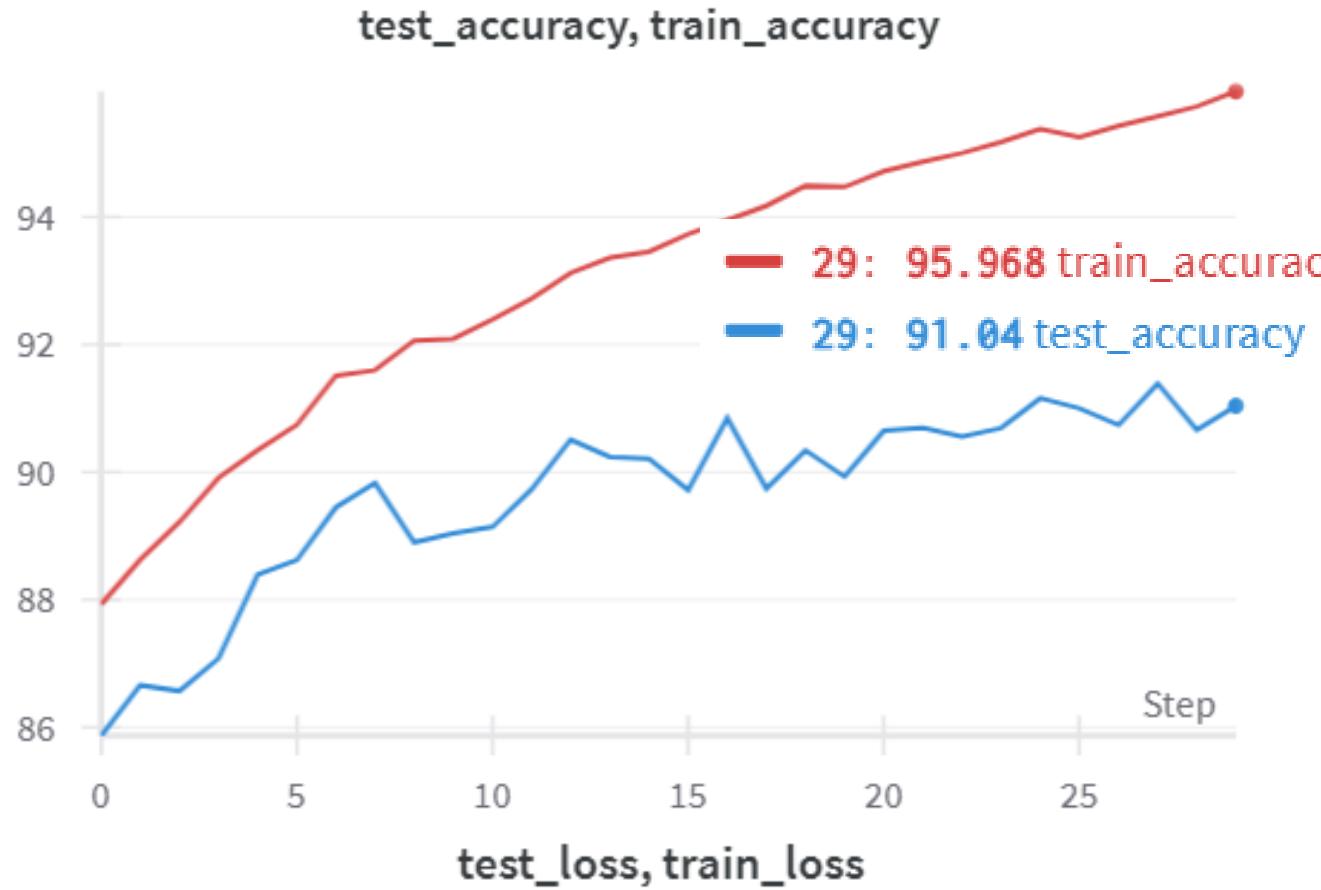


- Fine-tuning recovers most of the lost accuracy
- hard-ware efficiency with ps
- good accuracy with pu

Parameter	Value
ps	0.2
pu	0.2

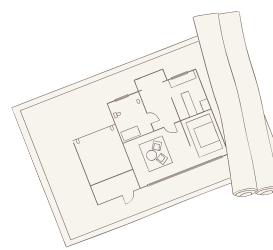


Pruning strategy



- higher ps
 - lower accuracy than ps=0.2
 - larger impact than unstructured alone

Parameter	Value
ps	0.3
pu	0.2



Factorization - Architecture changed

depthwise_from_scratch:

- modified version of the ResNet-18 model by replacing certain convolutional layers with depthwise separable convolutions, depending on the use_depthwise flag.

group: $2^{\text{puissance } x}$, where x in an integer

- modified version of the ResNet-18 model by replacing certain convolutional layers with depthwise separable convolutions, depending on the use_depthwise flag.

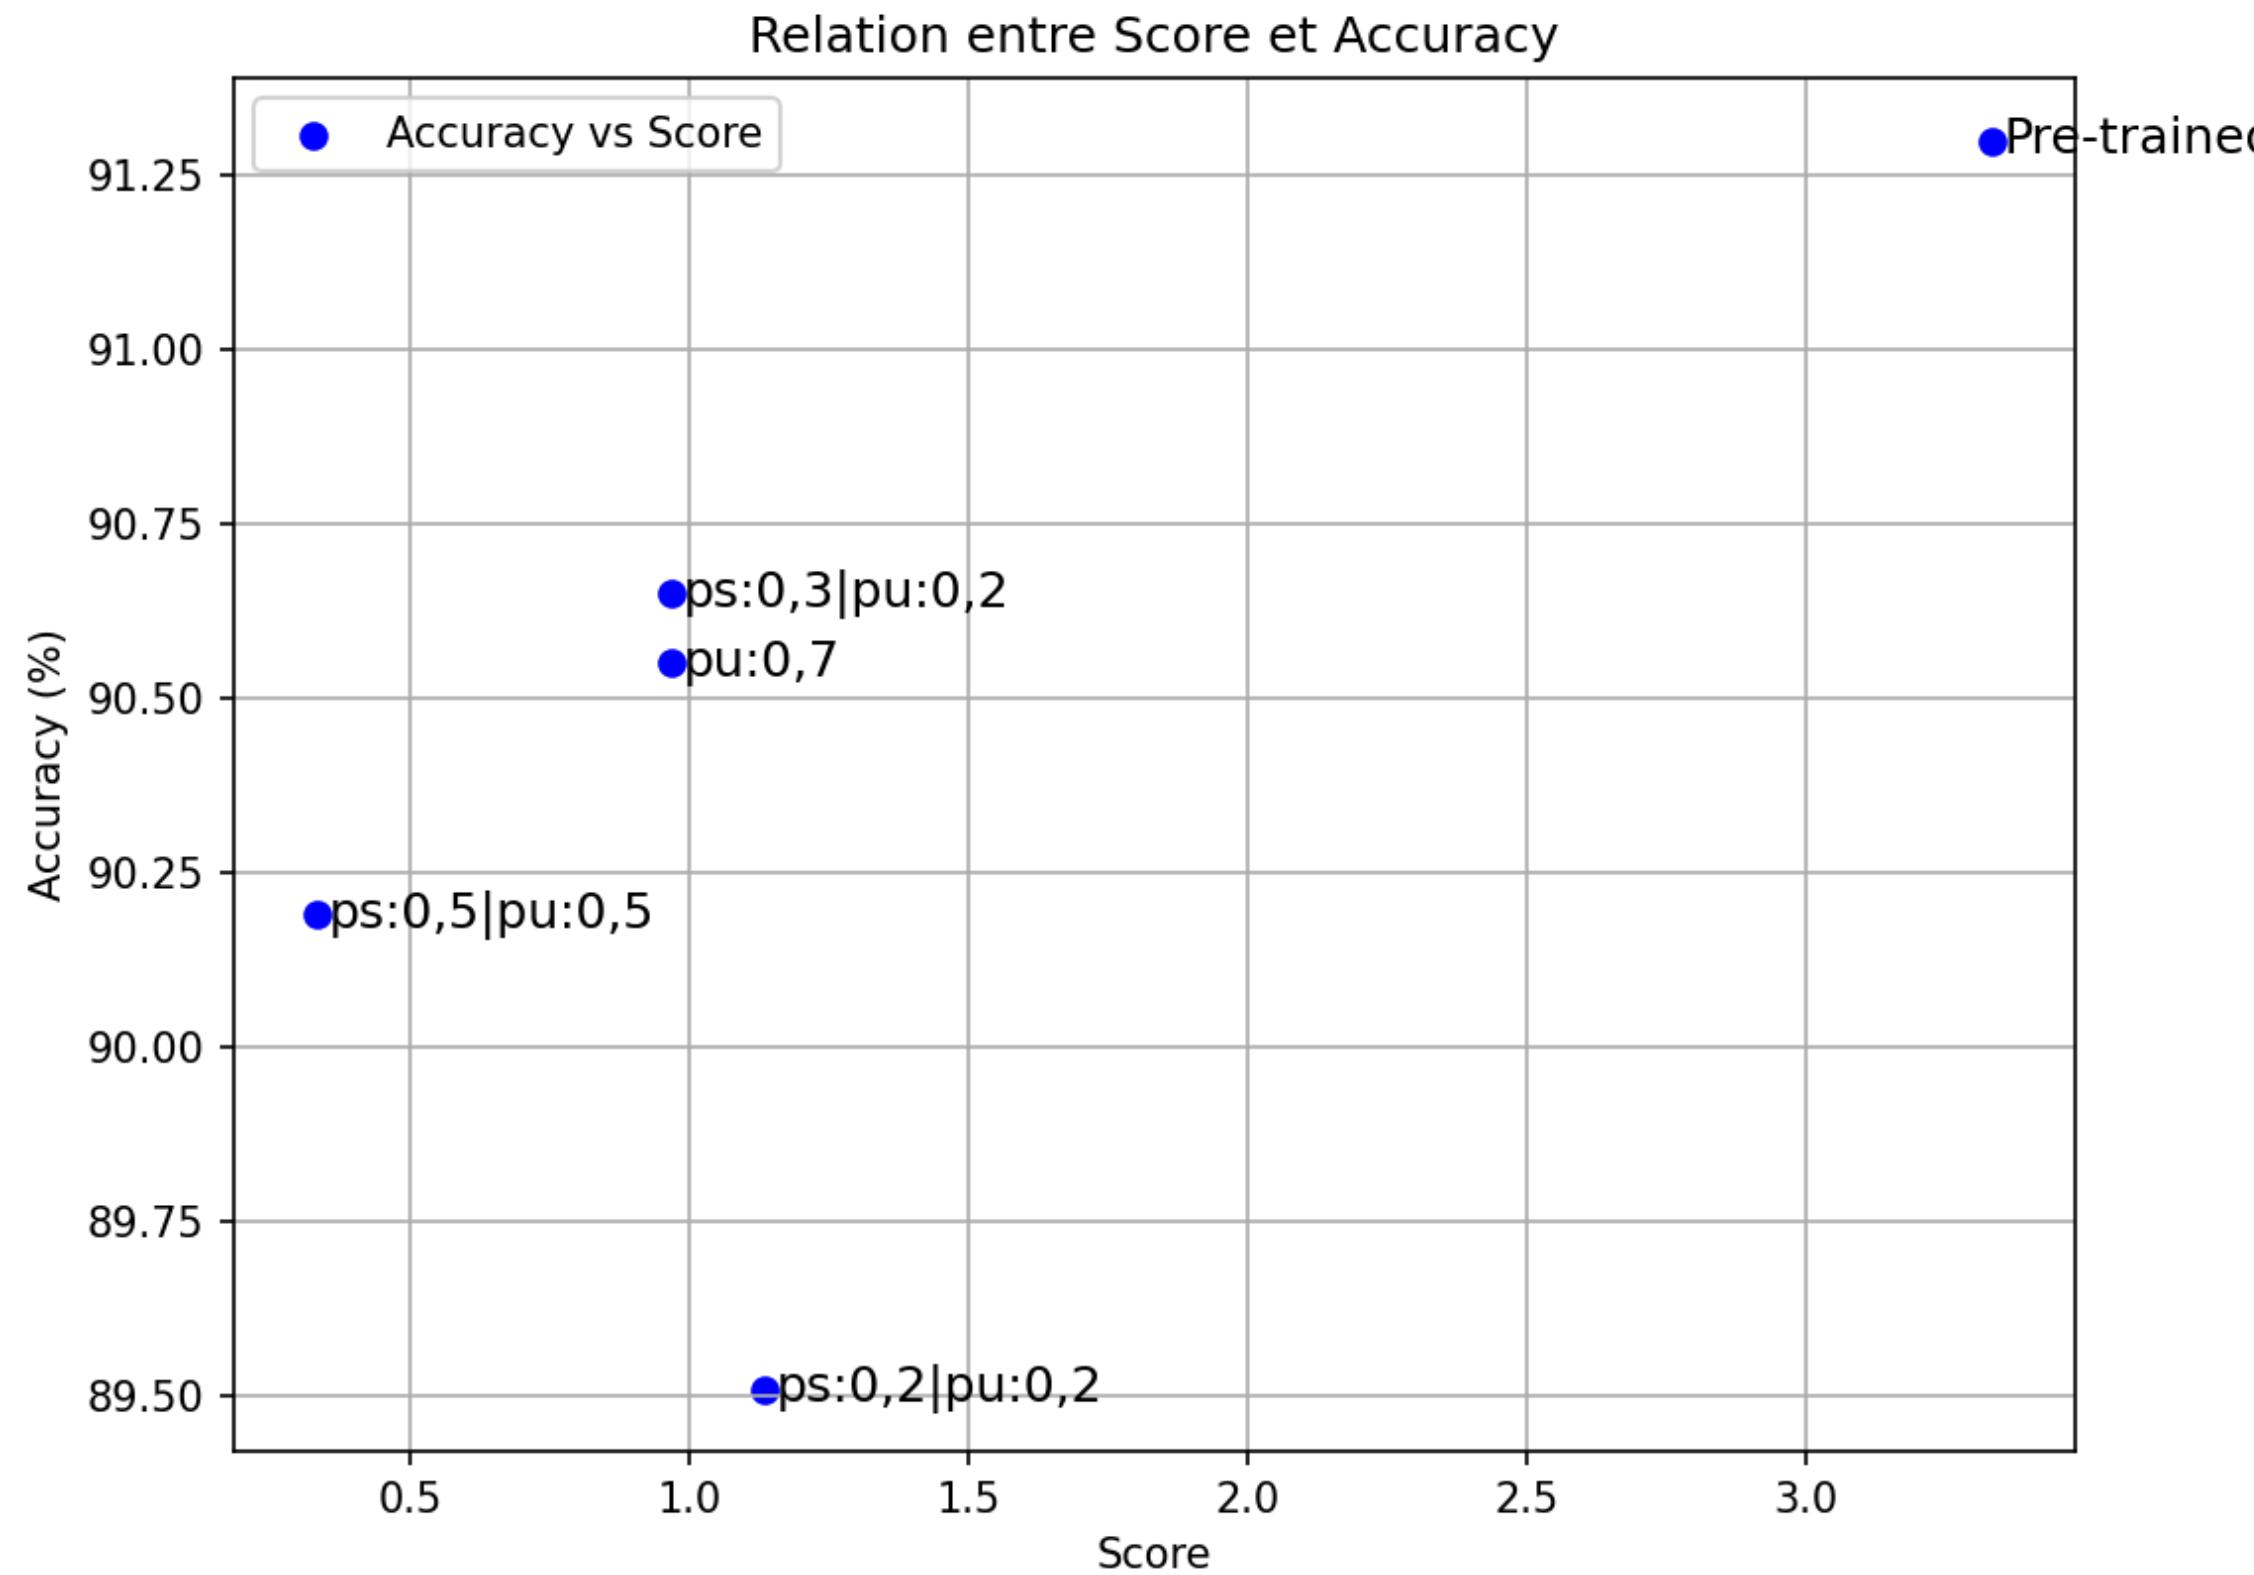
Score

$$\text{score} = \frac{[1 - (p_s + p_u)] \frac{q_w}{32} w}{\frac{5.6 \cdot 10^6}{\text{param}}} + \frac{(1 - p_s) \frac{\max(q_w, q_a)}{32} f}{\frac{2.8 \cdot 10^8}{\text{ops}}}$$

pre-trained model

Score	ps	pu	qa	qw
3.3337	-	-	-	-
0.9673	0.3	0.2	16	16
1.1340	0.2	0.2	16	16
0.9685	-	0.7	16	16

Conclusion



To do:

- Present different REsnet architecture
- Trade-off between score and accuracies
- Train depthwise on 100 epochs to see the evolution (can't track with wandb but can do an inference)